

Lab 3: Customer Segmentation

Machine Learning - MGT 665

Prof. Itauma

Group Members

- Prathyusha Devari
- Abdul Rehman
- Sai Srujana Jakkala
- Nirav

Introduction

For the project, we choose a Amazon customer data for analysis. Our task is to identify distinct customer groups based on their buying behaviors. Once segments are identified, analyze these groups to recommend specific marketing strategies tailored to each segment's characteristics.

Customer segmentation plays a vital role in helping businesses tailor their marketing strategies, promotions, and product offerings to distinct customer groups based on behaviors such as purchase frequency and spending patterns. Clustering techniques provide a data-driven way to identify such groups. In this report, we apply various clustering methods—K-Means, Hierarchical Clustering, DBSCAN, Agglomerative Clustering, and PCA (for visualization)—to a customer dataset. Each method brings unique strengths to segmenting customers and uncovering actionable insights.

Import necessary Libraries

```
import os
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from itertools import combinations
from sklearn.cluster import KMeans
from sklearn.decomposition import PCA
from scipy.stats import chi2_contingency
from sklearn.metrics import silhouette_score
from sklearn.preprocessing import StandardScaler

# To ignore warnings
import warnings
warnings.filterwarnings('ignore')
```

Dataset Overview

- age= age
- gender= gender

- Purchase_Frequency= How frequently do you make purchases on Amazon?
- Purchase_Categories= What product categories do you typically purchase on Amazon?
- Personalized_Recommendation_Frequency = Have you ever made a purchase based on personalized product recommendations from Amazon?
- Browsing_Frequency =How often do you browse Amazon's website or app?
- Product_Search_Method =How do you search for products on Amazon?
- Search_Result_Exploration =Do you tend to explore multiple pages of search results or focus on the first page?
- Customer_Reviews_Importance =How important are customer reviews in your decision-making process?
- Add_to_Cart_Browsing =Do you add products to your cart while browsing on Amazon?
- Cart_Completion_Frequency =How often do you complete the purchase after adding products to your cart?
- Cart_Abandonment_Factors =What factors influence your decision to abandon a purchase in your cart?
- Saveforlater_Frequency =Do you use Amazon's "Save for Later" feature, and if so, how often?
- Review_Left =Have you ever left a product review on Amazon?
- Review_Reliability =How much do you rely on product reviews when making a purchase?
- Review_Helpfulness =Do you find helpful information from other customers' reviews?
- Personalized_Recommendation_Frequency =How often do you receive personalized product recommendations from Amazon?
- Recommendation_Helpfulness =Do you find the recommendations helpful?
- Rating_Accuracy =How would you rate the relevance and accuracy of the recommendations you receive
- Shopping_Satisfaction =How satisfied are you with your overall shopping experience on Amazon?
- Service_Appreciation =What aspects of Amazon's services do you appreciate the most?
- Improvement_Areas =Are there any areas where you think Amazon can improve?

Reading the dataset

```
df = pd.read_csv('///Users//srujana//Downloads//Amazon Customer Behavior Survey.csv')
```

```
df.head()
```

	Timestamp	age	Gender	\
0	2023/06/04 1:28:19 PM GMT+5:30	23	Female	
1	2023/06/04 2:30:44 PM GMT+5:30	23	Female	
2	2023/06/04 5:04:56 PM GMT+5:30	24	Prefer not to say	
3	2023/06/04 5:13:00 PM GMT+5:30	24	Female	
4	2023/06/04 5:28:06 PM GMT+5:30	22	Female	

	Purchase_Frequency	Purchase_Categories	\
0	Few times a month	Beauty and	

Personal Care		
1	Once a month	Clothing and Fashion
2	Few times a month	Groceries and Gourmet Food;Clothing and Fashion
3	Once a month	Beauty and Personal Care;Clothing and Fashion;...
4	Less than once a month	Beauty and Personal Care;Clothing and Fashion

	Personalized_Recommendation_Frequency	Browsing_Frequency	\
0	Yes	Few times a week	
1	Yes	Few times a month	
2	No	Few times a month	
3	Sometimes	Few times a month	
4	Yes	Few times a month	

	Product_Search_Method	Search_Result_Exploration	\
0	Keyword	Multiple pages	
1	Keyword	Multiple pages	
2	Keyword	Multiple pages	
3	Keyword	First page	
4	Filter	Multiple pages	

	Customer_Reviews_Importance	...	Saveforlater_Frequency	Review_Left
0	1	...	Sometimes	Yes
1	1	...	Rarely	No
2	2	...	Rarely	No
3	5	...	Sometimes	Yes
4	1	...	Rarely	No

	Review_Reliability	Review_Helpfulness	\
0	Occasionally	Yes	
1	Heavily	Yes	
2	Occasionally	No	
3	Heavily	Yes	
4	Heavily	Yes	

	Personalized_Recommendation_Frequency	Recommendation_Helpfulness	\
0	2	Yes	
1	2	Sometimes	
2	4	No	
3	3	Sometimes	
4	4	Yes	

	Rating_Accuracy	Shopping_Satisfaction	Service_Appreciation \
0	1	1	Competitive prices
1	3	2	Wide product selection
2	3	3	Competitive prices
3	3	4	Competitive prices
4	2	2	Competitive prices

	Improvement_Areas
0	Reducing packaging waste
1	Reducing packaging waste
2	Product quality and accuracy
3	Product quality and accuracy
4	Product quality and accuracy

[5 rows x 23 columns]

df.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 602 entries, 0 to 601

Data columns (total 23 columns):

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	Timestamp	602 non-null	object
1	age	602 non-null	int64
2	Gender	602 non-null	object
3	Purchase_Frequency	602 non-null	object
4	Purchase_Categories	602 non-null	object
5	Personalized_Recommendation_Frequency	602 non-null	object
6	Browsing_Frequency	602 non-null	object
7	Product_Search_Method	600 non-null	object
8	Search_Result_Exploration	602 non-null	object
9	Customer_Reviews_Importance	602 non-null	int64
10	Add_to_Cart_Browsing	602 non-null	object
11	Cart_Completion_Frequency	602 non-null	object
12	Cart_Abandonment_Factors	602 non-null	object
13	Saveforlater_Frequency	602 non-null	object
14	Review_Left	602 non-null	object
15	Review_Reliability	602 non-null	object
16	Review_Helpfulness	602 non-null	object
17	Personalized_Recommendation_Frequency	602 non-null	int64
18	Recommendation_Helpfulness	602 non-null	object
19	Rating_Accuracy	602 non-null	int64
20	Shopping_Satisfaction	602 non-null	int64
21	Service_Appreciation	602 non-null	object
22	Improvement_Areas	602 non-null	object

dtypes: int64(5), object(18)

memory usage: 108.3+ KB

There are only 2 datatypes int and object

Null value check

```
df.isna().sum()

Timestamp                                0
age                                       0
Gender                                   0
Purchase_Frequency                      0
Purchase_Categories                     0
Personalized_Recommendation_Frequency  0
Browsing_Frequency                     0
Product_Search_Method                   2
Search_Result_Exploration               0
Customer_Reviews_Importance             0
Add_to_Cart_Browsing                    0
Cart_Completion_Frequency               0
Cart_Abandonment_Factors                0
Saveforlater_Frequency                 0
Review_Left                             0
Review_Reliability                     0
Review_Helpfulness                     0
Personalized_Recommendation_Frequency  0
Recommendation_Helpfulness              0
Rating_Accuracy                        0
Shopping_Satisfaction                   0
Service_Appreciation                   0
Improvement_Areas                      0
dtype: int64
```

Only the Product_Search_Method has null values and this could be ignored for our analysis

Catagorie the numerical and catagorical data

```
# Numerical columns
numerical = df.select_dtypes(include=['int64'])

# Categorical columns
categorical = df.select_dtypes(include=['object'])

# Total numerical data
numerical

   age  Customer_Reviews_Importance
0   23                             1
2
1   23                             1
2
```

```

2      24      2
4
3      24      5
3
4      22      1
4
..      ...      ...
...
597    23      4
3
598    23      3
3
599    23      3
3
600    23      1
2
601    23      3
3

```

	Rating_Accuracy	Shopping_Satisfaction
0	1	1
1	3	2
2	3	3
3	3	4
4	2	2
..
597	3	4
598	3	3
599	2	3
600	2	2
601	3	3

```
[602 rows x 5 columns]
```

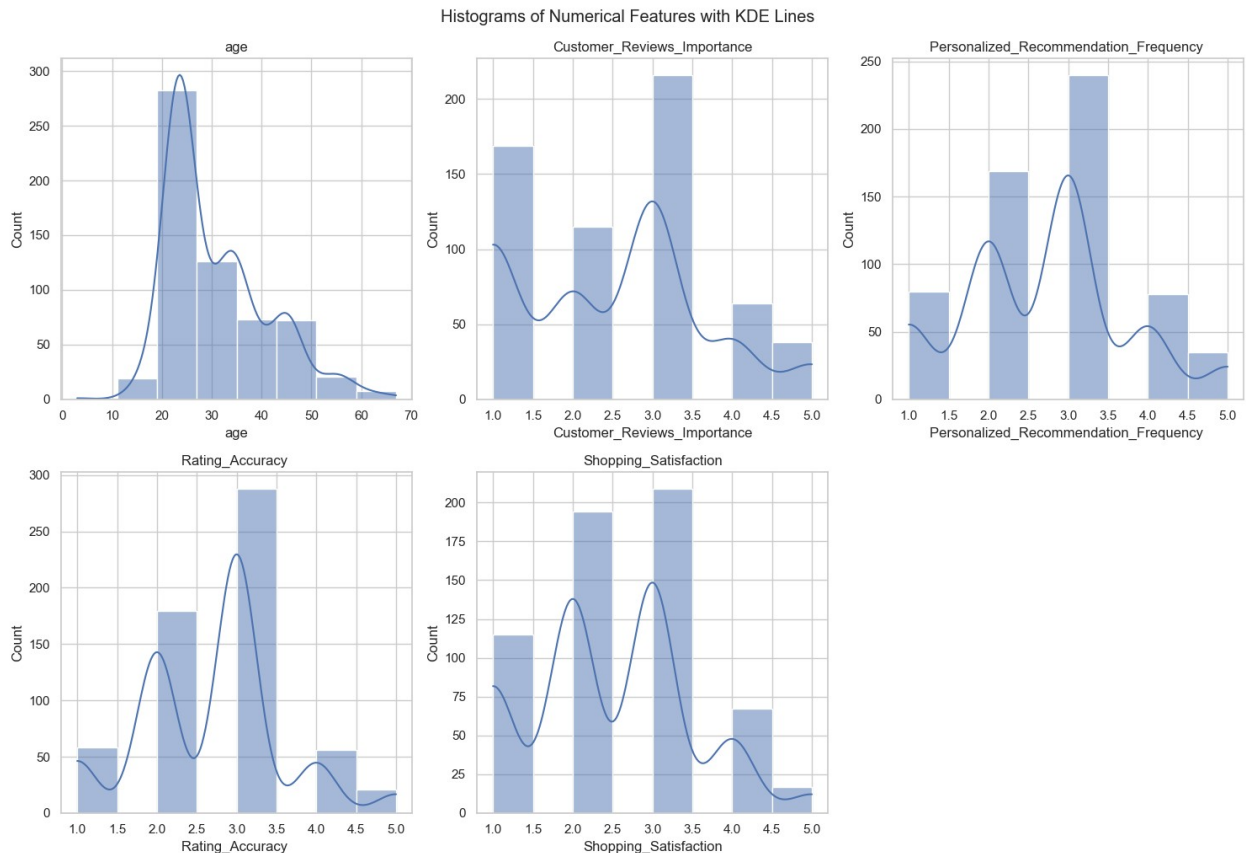
Methodology

Exploratory Data Analysis

```

# Histograms with KDE lines
sns.set(style="whitegrid")
plt.figure(figsize=(15, 10))
for i, column in enumerate(numerical.columns, 1):
    plt.subplot(2, 3, i)
    sns.histplot(numerical[column], kde=True, bins=8)
    plt.title(f'{column}')
plt.tight_layout()
plt.suptitle('Histograms of Numerical Features with KDE Lines',
y=1.02)
plt.show()

```

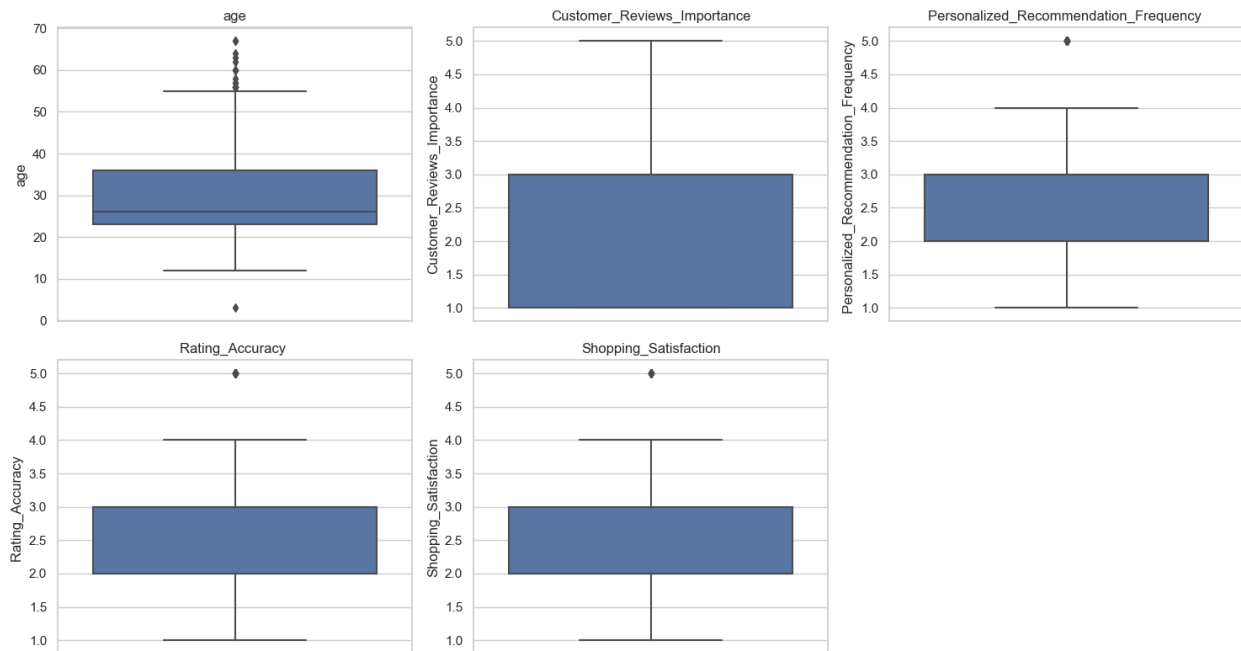


From the graphs:

- The age has a positive skew, it means that there are a few high-value outliers (besides that 3 on the left) on the right side
- Customer_Review_Importance has a significant concentration in number 1 and 3 turning the distribution not symmetric and reflecting the bimodal, thus causing the mean to drop in comparison to the rest where and. This suggests that there are two main groups of customers: group that does not find customer reviews important (rating 1) and group finds them moderately important (rating 3)
- Personalized_Recommendation_Frequency shows a distribution with a clear peak at rating 3, suggesting that most customers experience a moderate frequency of personalized recommendations and there is a slight left skew (negative skew), with a tail extending towards the lower ratings
- Rating_Accuracy shows a peak at rating 3, indicating that most customers rate the accuracy of ratings as neutral or moderate and there is a slight left skew (negative skew)
- Shopping_Satisfaction appears to have a peak at rating 3, with a significant number of responses also at rating 2 indicating that a number of customers are less satisfied and a pronounced concentration in the 3 first quarters

Outliers Check

```
# Box Plot
plt.figure(figsize=(15, 8))
for i, column in enumerate(numerical.columns, 1):
    plt.subplot(2, 3, i)
    sns.boxplot(y=numerical[column])
    plt.title(f'{column}')
plt.tight_layout()
plt.show()
```




From the boxplots:

- age has this line indicating the median age (mid-20s), and there are several outliers on both the lower and upper ends, with more outliers on the upper end, indicating some customers are significantly older than the average.
- Customer_Reviews_Importance box is compact, indicating that the middle 50% of ratings are close together and suggesting that some customers rate the importance of customer reviews as significantly lower than the rest.
- Personalized_Recommendation_Frequency, Rating_Accuracy and Shopping_Satisfaction box appears to be very short, showing little variation in the middle 50% and has a few outliers on the upper end.

```
# Correlation Matrix
correlation_matrix = numerical.corr()
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm',
```



```
linewidths=0.5)
plt.title('Correlation Matrix')
plt.show()
```



EDA for Catagorical data

categorical

	Timestamp				Gender \
0	2023/06/04	1:28:19	PM	GMT+5:30	Female
1	2023/06/04	2:30:44	PM	GMT+5:30	Female
2	2023/06/04	5:04:56	PM	GMT+5:30	Prefer not to say
3	2023/06/04	5:13:00	PM	GMT+5:30	Female
4	2023/06/04	5:28:06	PM	GMT+5:30	Female
..

597	2023/06/12	4:02:02	PM	GMT+5:30	Female
598	2023/06/12	4:02:53	PM	GMT+5:30	Female
599	2023/06/12	4:03:59	PM	GMT+5:30	Female
600	2023/06/12	9:57:20	PM	GMT+5:30	Female
601	2023/06/16	9:16:05	AM	GMT+5:30	Female

Purchase_Frequency \	
0	Few times a month
1	Once a month
2	Few times a month
3	Once a month
4	Less than once a month
..	...
597	Once a week
598	Once a week
599	Once a month
600	Few times a month
601	Once a week

Purchase_Categories \	
0	Beauty and Personal Care
1	Clothing and Fashion
2	Groceries and Gourmet Food;Clothing and Fashion
3	Beauty and Personal Care;Clothing and Fashion;...
4	Beauty and Personal Care;Clothing and Fashion
..	...
597	Beauty and Personal Care
598	Clothing and Fashion
599	Beauty and Personal Care
600	Beauty and Personal Care;Clothing and Fashion;...
601	Clothing and Fashion

Personalized_Recommendation_Frequency	Browsing_Frequency \
0	Yes Few times a week
1	Yes Few times a month
2	No Few times a month
3	Sometimes Few times a month
4	Yes Few times a month
..	...
597	Sometimes Few times a week
598	Sometimes Few times a week
599	Sometimes Few times a week
600	Yes Few times a month
601	Sometimes Multiple times a day

Product_Search_Method	Search_Result_Exploration
Add_to_Cart_Browsing \	
0	Keyword Multiple pages
Yes	
1	Keyword Multiple pages

Yes			
2	Keyword	Multiple pages	
Yes			
3	Keyword	First page	
Maybe			
4	Filter	Multiple pages	
Yes			
..
..			
597	categories	Multiple pages	
Maybe			
598	Filter	Multiple pages	
Maybe			
599	categories	Multiple pages	
Maybe			
600	Keyword	Multiple pages	
Yes			
601	Keyword	Multiple pages	
Maybe			
	Cart_Completion_Frequency	Cart_Abandonment_Factors	\
0	Sometimes	Found a better price elsewhere	
1	Often	High shipping costs	
2	Sometimes	Found a better price elsewhere	
3	Sometimes	Found a better price elsewhere	
4	Sometimes	High shipping costs	
..	
597	Sometimes	Found a better price elsewhere	
598	Sometimes	Found a better price elsewhere	
599	Sometimes	High shipping costs	
600	Often	others	
601	Often	Found a better price elsewhere	
	Saveforlater_Frequency	Review_Left	Review_Reliability
	Review_Helpfulness	\	
0	Sometimes	Yes	Occasionally
Yes			
1	Rarely	No	Heavily
Yes			
2	Rarely	No	Occasionally
No			
3	Sometimes	Yes	Heavily
Yes			
4	Rarely	No	Heavily
Yes			
..
...			
597	Sometimes	Yes	Moderately
Sometimes			

598	Sometimes	Yes	Heavily
Sometimes			
599	Sometimes	Yes	Occasionally
Sometimes			
600	Sometimes	No	Heavily
Yes			
601	Sometimes	Yes	Moderately
Sometimes			

	Recommendation_Helpfulness	Service_Appreciation \
0	Yes	Competitive prices
1	Sometimes	Wide product selection
2	No	Competitive prices
3	Sometimes	Competitive prices
4	Yes	Competitive prices
..
597	Sometimes	Competitive prices
598	Sometimes	Product recommendations
599	Sometimes	Wide product selection
600	Yes	Wide product selection
601	Sometimes	Product recommendations

	Improvement_Areas
0	Reducing packaging waste
1	Reducing packaging waste
2	Product quality and accuracy
3	Product quality and accuracy
4	Product quality and accuracy
..	...
597	Customer service responsiveness
598	Reducing packaging waste
599	Product quality and accuracy
600	Product quality and accuracy
601	Product quality and accuracy

[602 rows x 18 columns]

Convert Timestamp to datetime

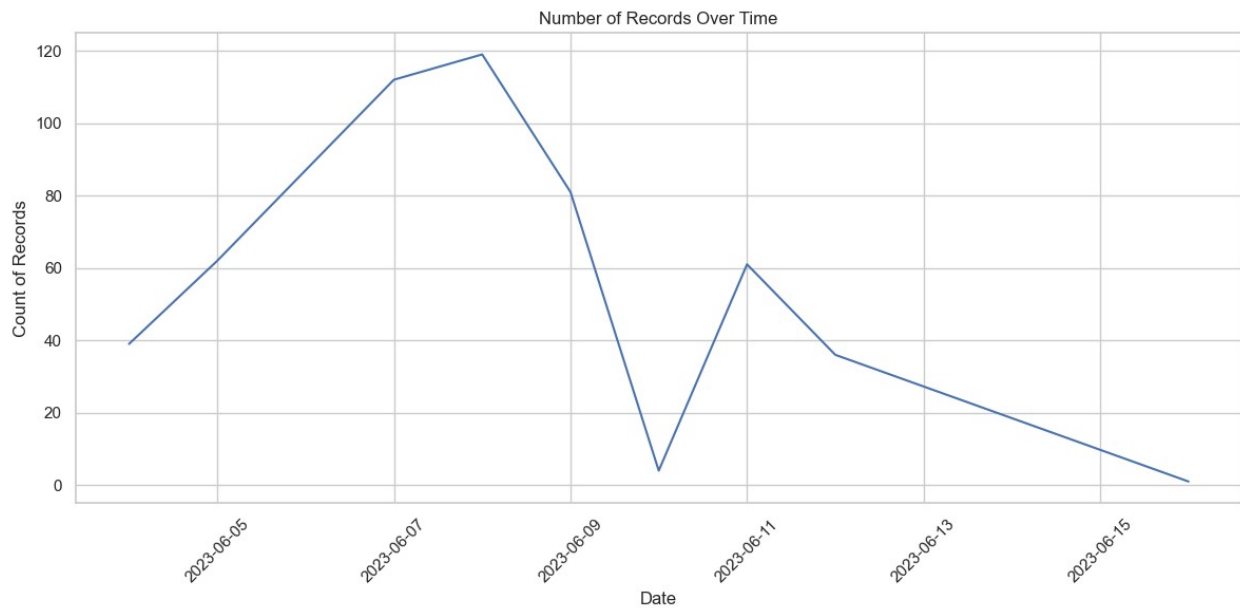
```
categorical['Timestamp'] = pd.to_datetime(categorical['Timestamp'])
time_series_data =
categorical['Timestamp'].dt.date.value_counts().sort_index()
```

```
sns.set(style="whitegrid")
```

Line Plot

```
plt.figure(figsize=(12, 6))
sns.lineplot(x=time_series_data.index, y=time_series_data.values)
plt.title('Number of Records Over Time')
plt.xlabel('Date')
plt.ylabel('Count of Records')
```

```
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



```
print("Starts: ", categorical['Timestamp'].min())
print("Ends: ", categorical['Timestamp'].max())
```

```
Starts: 2023-06-04 13:28:19-05:30
Ends: 2023-06-16 09:16:05-05:30
```

```
def visualize_categorical(series, title=None):
    # Value counts
    print(series.value_counts())

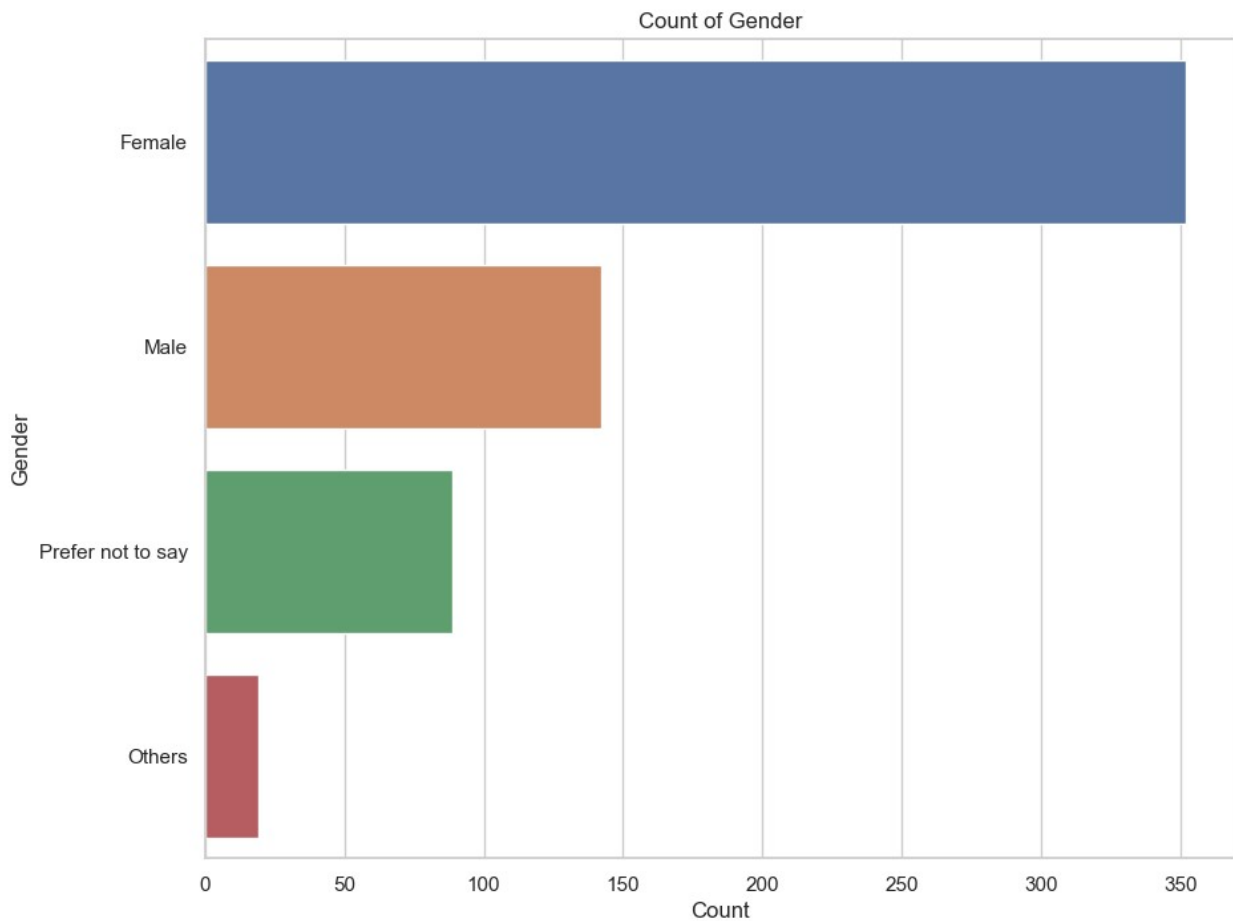
    # Set up the seaborn style
    sns.set(style="whitegrid")

    # Count Plot
    plt.figure(figsize=(10, 8))
    sns.countplot(y=series, order=series.value_counts().index)
    plt.title(title if title else f'Count of {series.name}')
    plt.xlabel('Count')
    plt.ylabel(series.name)
    plt.show()
```

```
visualize_categorical(categorical['Gender'])
```

```
Gender
Female    352
Male     142
```

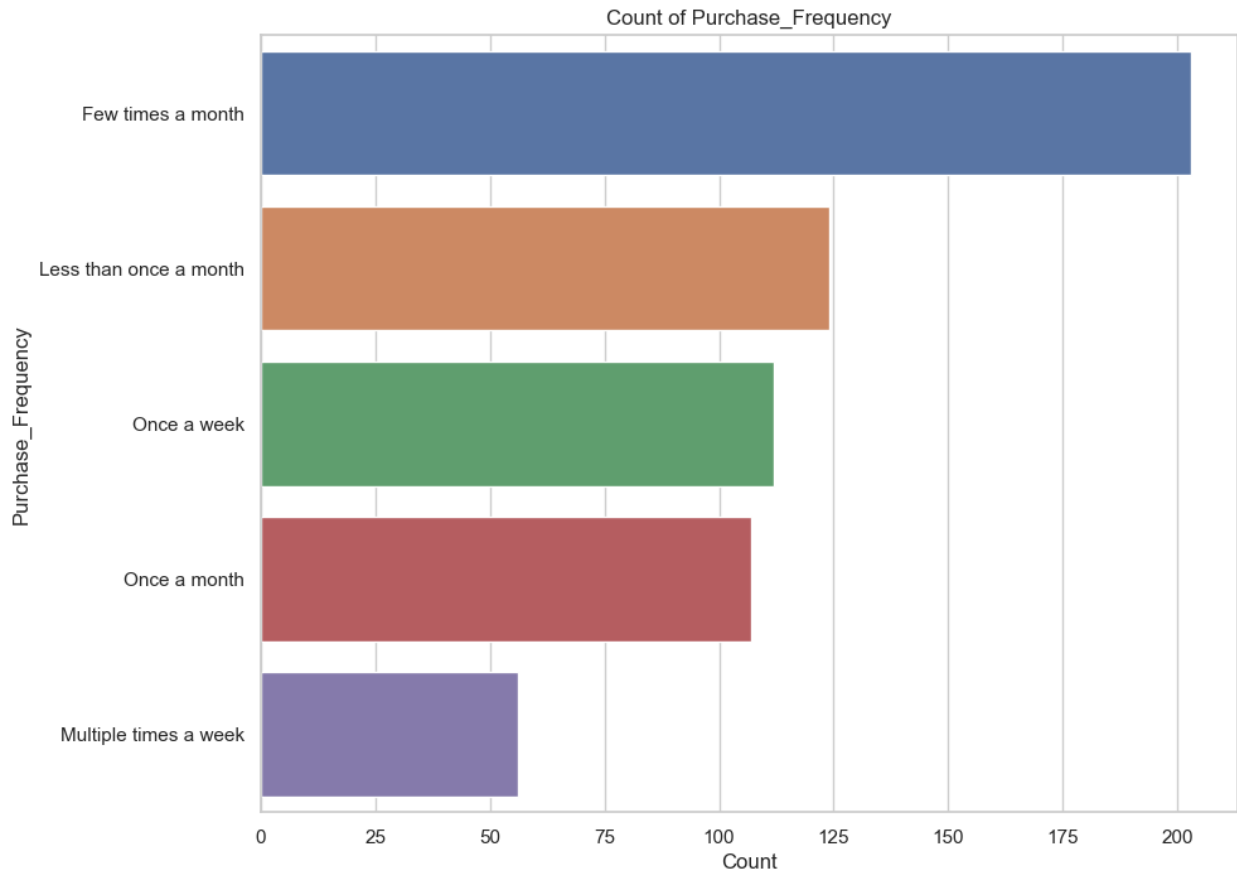
```
Prefer not to say    89
Others               19
Name: count, dtype: int64
```



A significantly higher number of females are represented in this dataset compared to males and other gender categories

```
visualize_categorical(categorical['Purchase_Frequency'])
```

```
Purchase_Frequency
Few times a month      203
Less than once a month 124
Once a week           112
Once a month          107
Multiple times a week   56
Name: count, dtype: int64
```



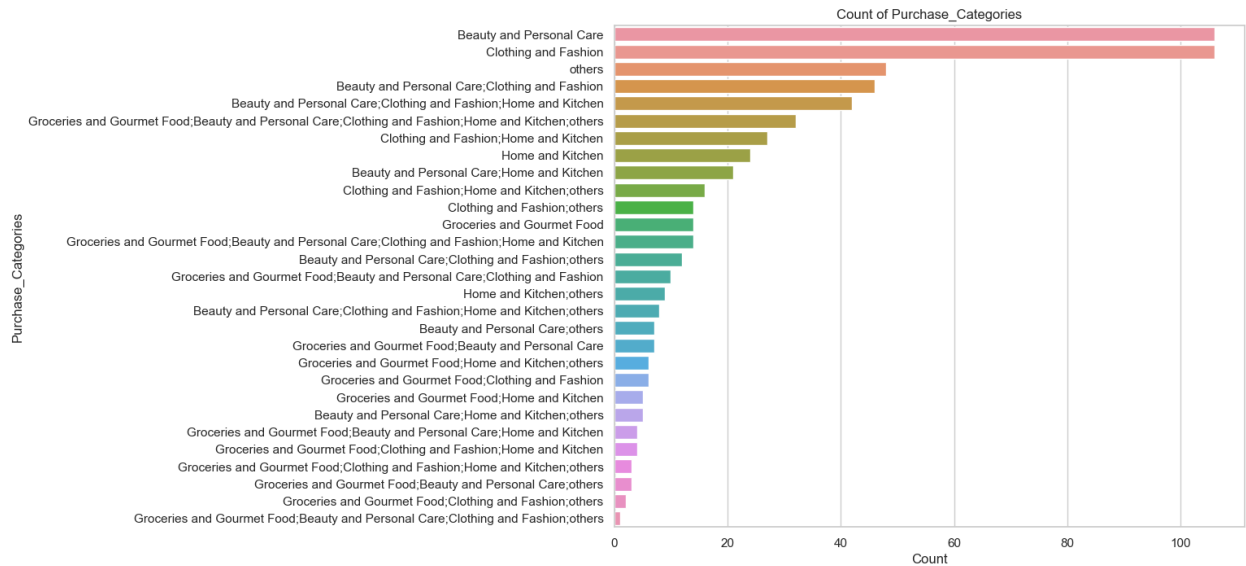
The majority of customers make purchases a few times a month, followed by those purchasing less than once a month. This suggests that most customers are occasional shoppers rather than frequent ones

```
visualize_categorical(categorical['Purchase_Categories'])
Purchase_Categories
Beauty and Personal Care
106
Clothing and Fashion
106
others
48
Beauty and Personal Care;Clothing and Fashion
46
Beauty and Personal Care;Clothing and Fashion;Home and Kitchen
42
Groceries and Gourmet Food;Beauty and Personal Care;Clothing and
Fashion;Home and Kitchen;others    32
Clothing and Fashion;Home and Kitchen
27
Home and Kitchen
24
```

```

Beauty and Personal Care;Home and Kitchen
21
Clothing and Fashion;Home and Kitchen;others
16
Clothing and Fashion;others
14
Groceries and Gourmet Food
14
Groceries and Gourmet Food;Beauty and Personal Care;Clothing and
Fashion;Home and Kitchen 14
Beauty and Personal Care;Clothing and Fashion;others
12
Groceries and Gourmet Food;Beauty and Personal Care;Clothing and
Fashion 10
Home and Kitchen;others
9
Beauty and Personal Care;Clothing and Fashion;Home and Kitchen;others
8
Beauty and Personal Care;others
7
Groceries and Gourmet Food;Beauty and Personal Care
7
Groceries and Gourmet Food;Home and Kitchen;others
6
Groceries and Gourmet Food;Clothing and Fashion
6
Groceries and Gourmet Food;Home and Kitchen
5
Beauty and Personal Care;Home and Kitchen;others
5
Groceries and Gourmet Food;Beauty and Personal Care;Home and Kitchen
4
Groceries and Gourmet Food;Clothing and Fashion;Home and Kitchen
4
Groceries and Gourmet Food;Clothing and Fashion;Home and
Kitchen;others 3
Groceries and Gourmet Food;Beauty and Personal Care;others
3
Groceries and Gourmet Food;Clothing and Fashion;others
2
Groceries and Gourmet Food;Beauty and Personal Care;Clothing and
Fashion;others 1
Name: count, dtype: int64

```

The most frequent product categories purchased are Beauty, Personal Care, Clothing and Fashion.

```
visualize_categorical(categorical['Personalized_Recommendation_Frequency'])
```

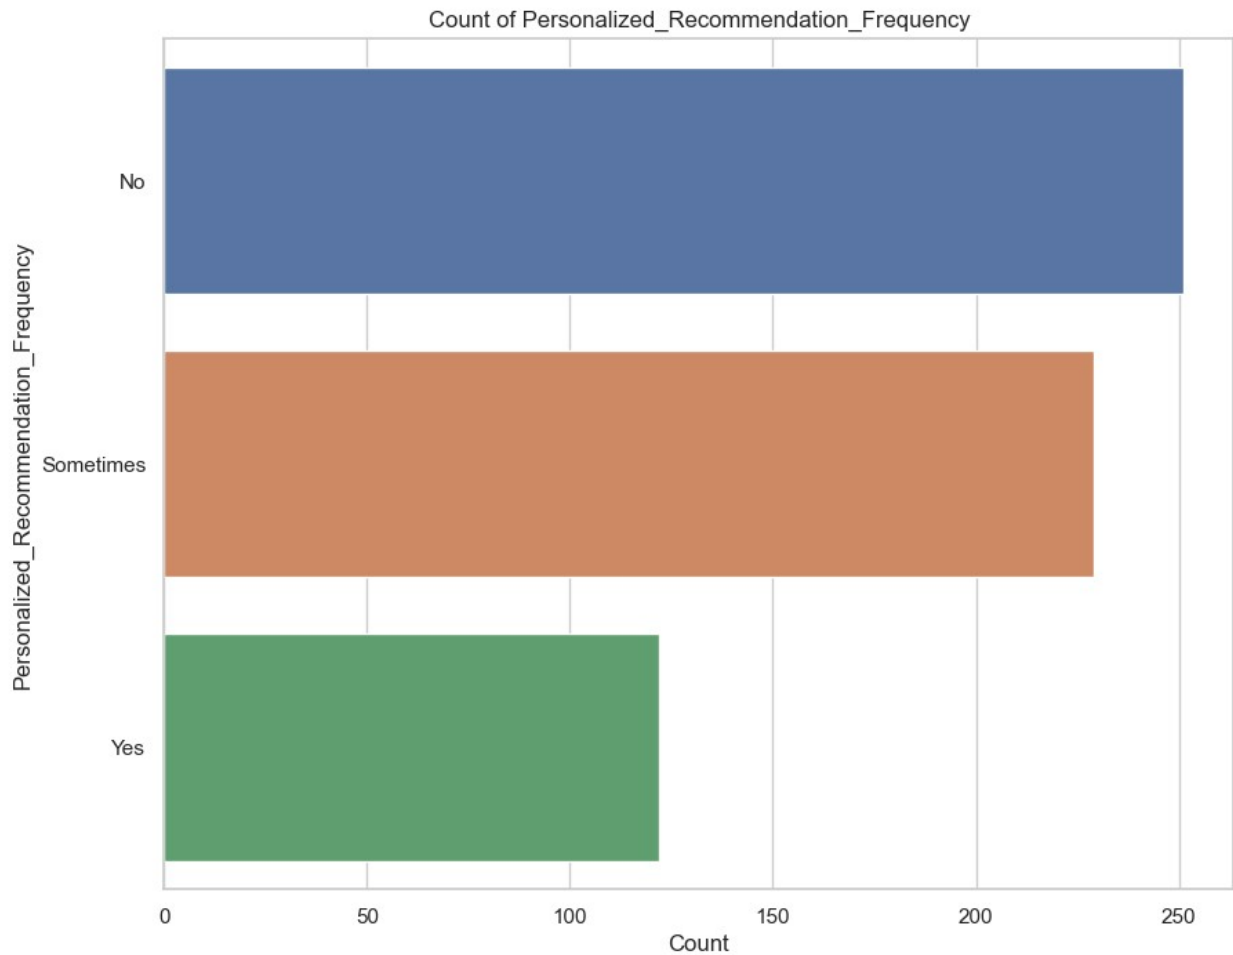
Personalized_Recommendation_Frequency

No 251

Sometimes 229

Yes 122

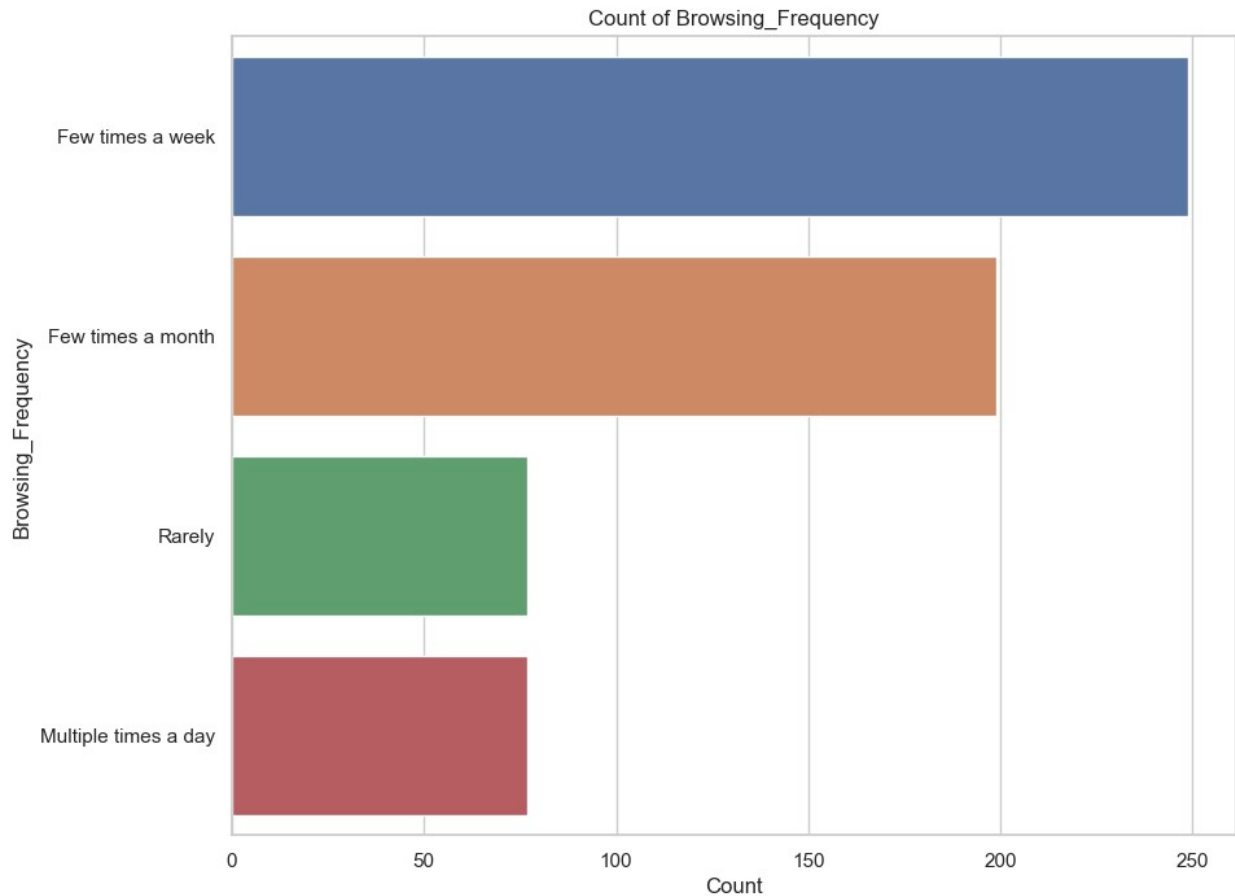
Name: count, dtype: int64



Many customers answered with 'No' being the most common response, suggesting a significant number of customers may not be influenced by or aware of personalized recommendations

```
visualize_categorical(categorical['Browsing_Frequency'])
```

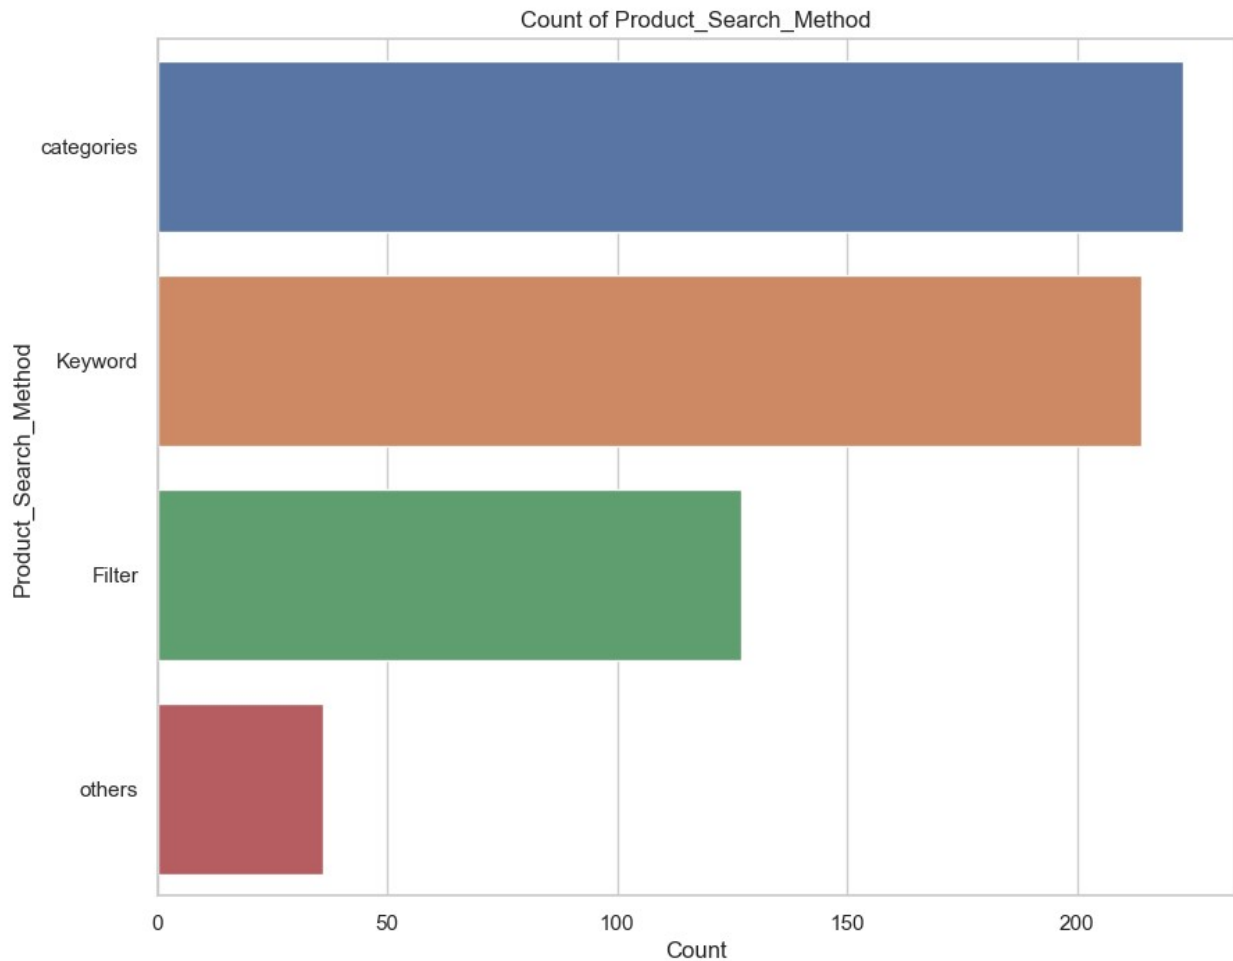
```
Browsing_Frequency
Few times a week      249
Few times a month     199
Rarely                77
Multiple times a day   77
Name: count, dtype: int64
```



Few times a week is the most common response, indicating regular engagement with the Amazon platform, though not necessarily daily

```
visualize_categorical(categorical['Product_Search_Method'])
```


```
Product_Search_Method
categories      223
Keyword        214
Filter          127
others          36
Name: count, dtype: int64
```



The most common search method is by categories, followed by keyword searches. Fewer users utilize filters or other unspecified methods, suggesting that users prefer broad search methods, possibly to discover a wider range of options.

```
visualize_categorical(categorical['Search_Result_Exploration'])
```

```
Search_Result_Exploration
Multiple pages    442
First page       160
Name: count, dtype: int64
```



The majority of users tend to explore multiple pages of search results rather than only the first page, suggesting that users are looking for more options before making a decision

```
visualize_categorical(categorical['Add_to_Cart_Browsing'])
```

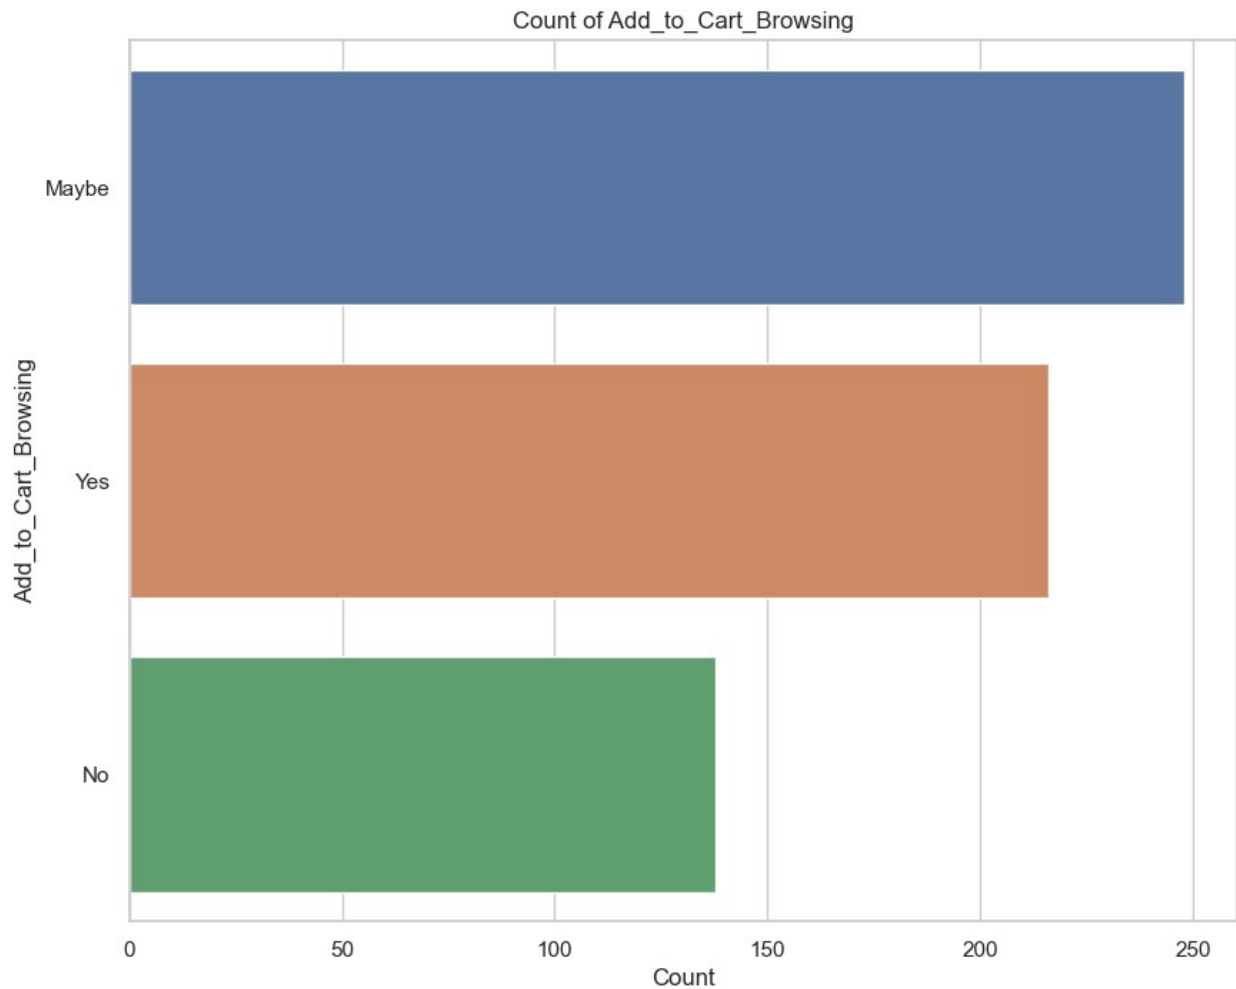
```
Add_to_Cart_Browsing
```

```
Maybe    248
```

```
Yes       216
```

```
No        138
```

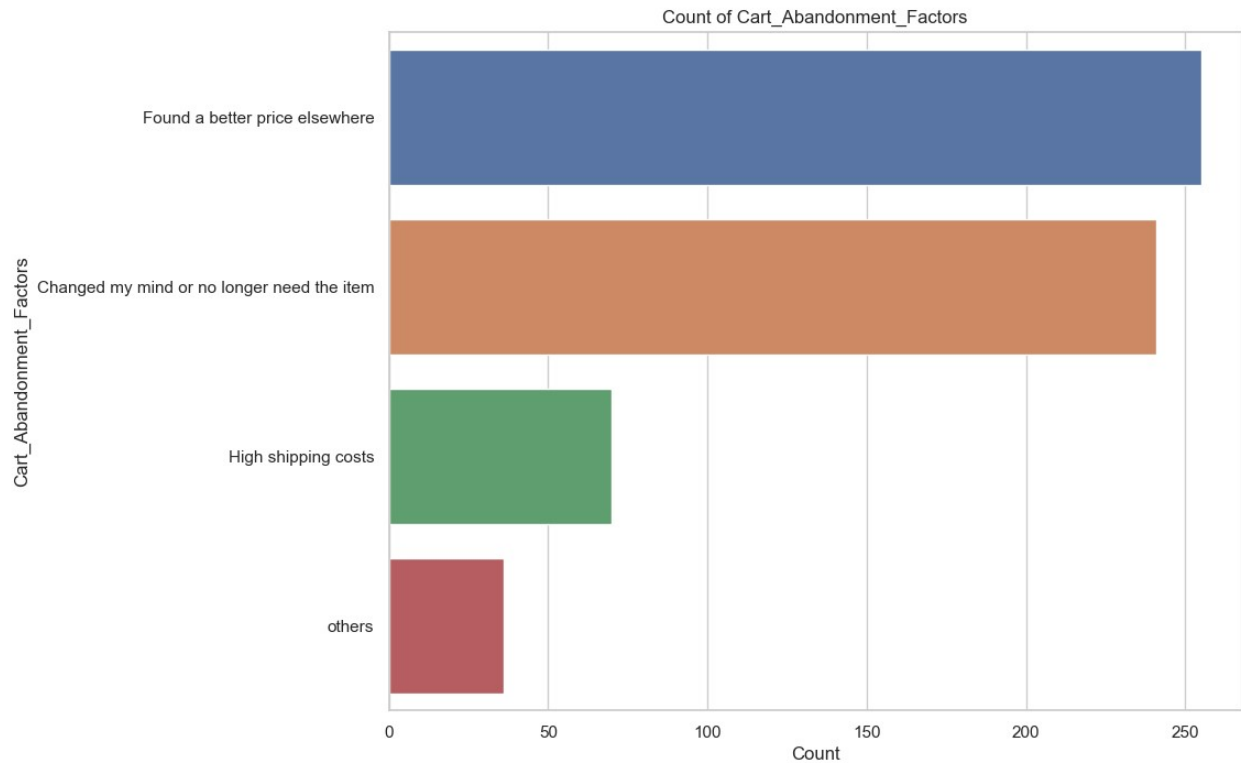
```
Name: count, dtype: int64
```



A significant number of customers add items to their cart while browsing, but 'Maybe' is the most common response, suggesting that customers are selective about what they add to their cart

```
visualize_categorical(categorical['Cart_Completion_Frequency'])
```

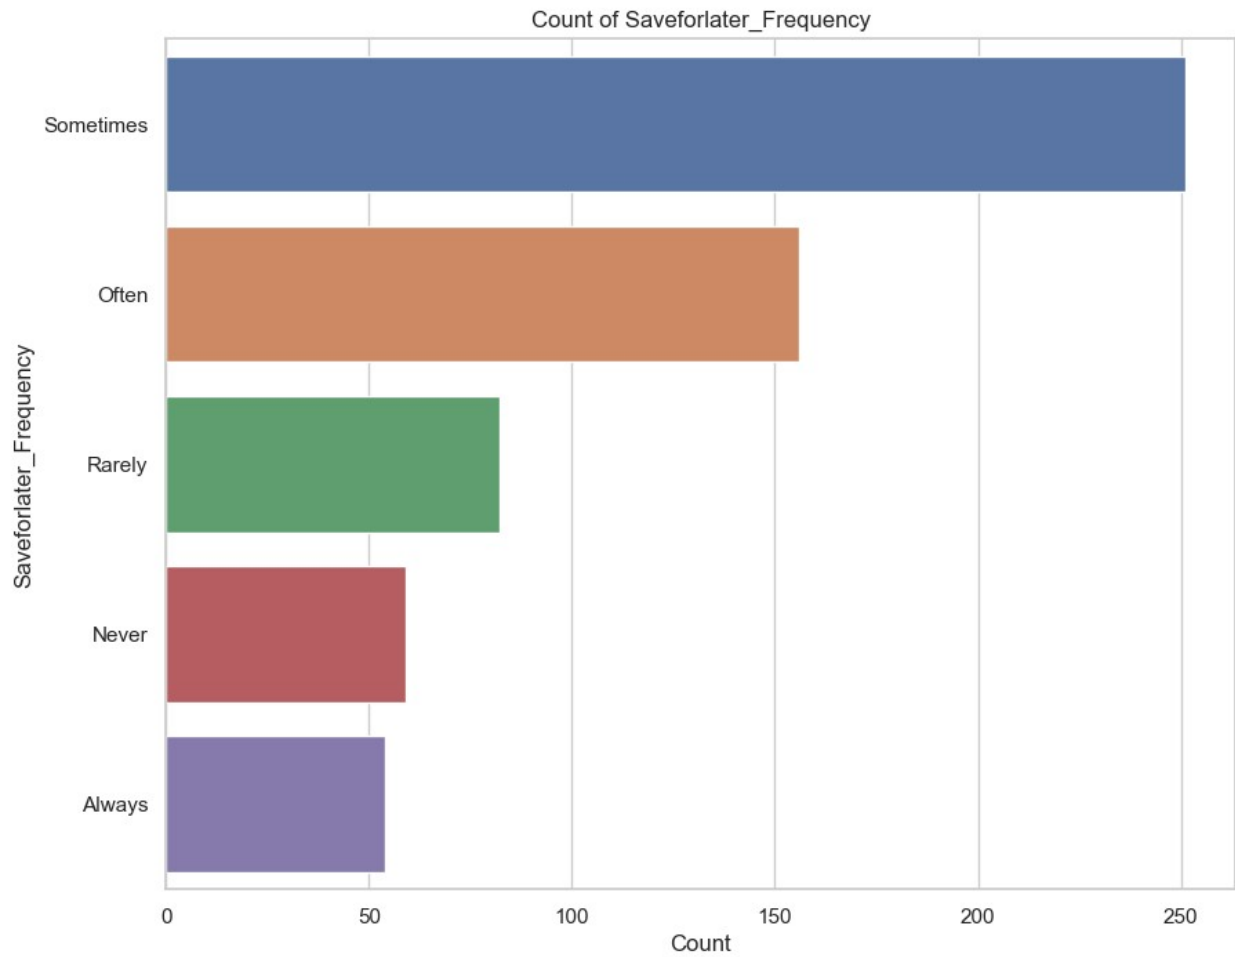
```
Cart_Completion_Frequency
Sometimes      304
Often         158
Rarely         72
Always         47
Never          21
Name: count, dtype: int64
```

The top reasons for cart abandonment are Found a better price elsewhere and Changed my mind or no longer need the item, highlighting price sensitivity and changing customer needs as key factors

```
visualize_categorical(categorical['Saveforlater_Frequency'])
```

```
Saveforlater_Frequency
Sometimes      251
Often         156
Rarely         82
Never          59
Always         54
Name: count, dtype: int64
```

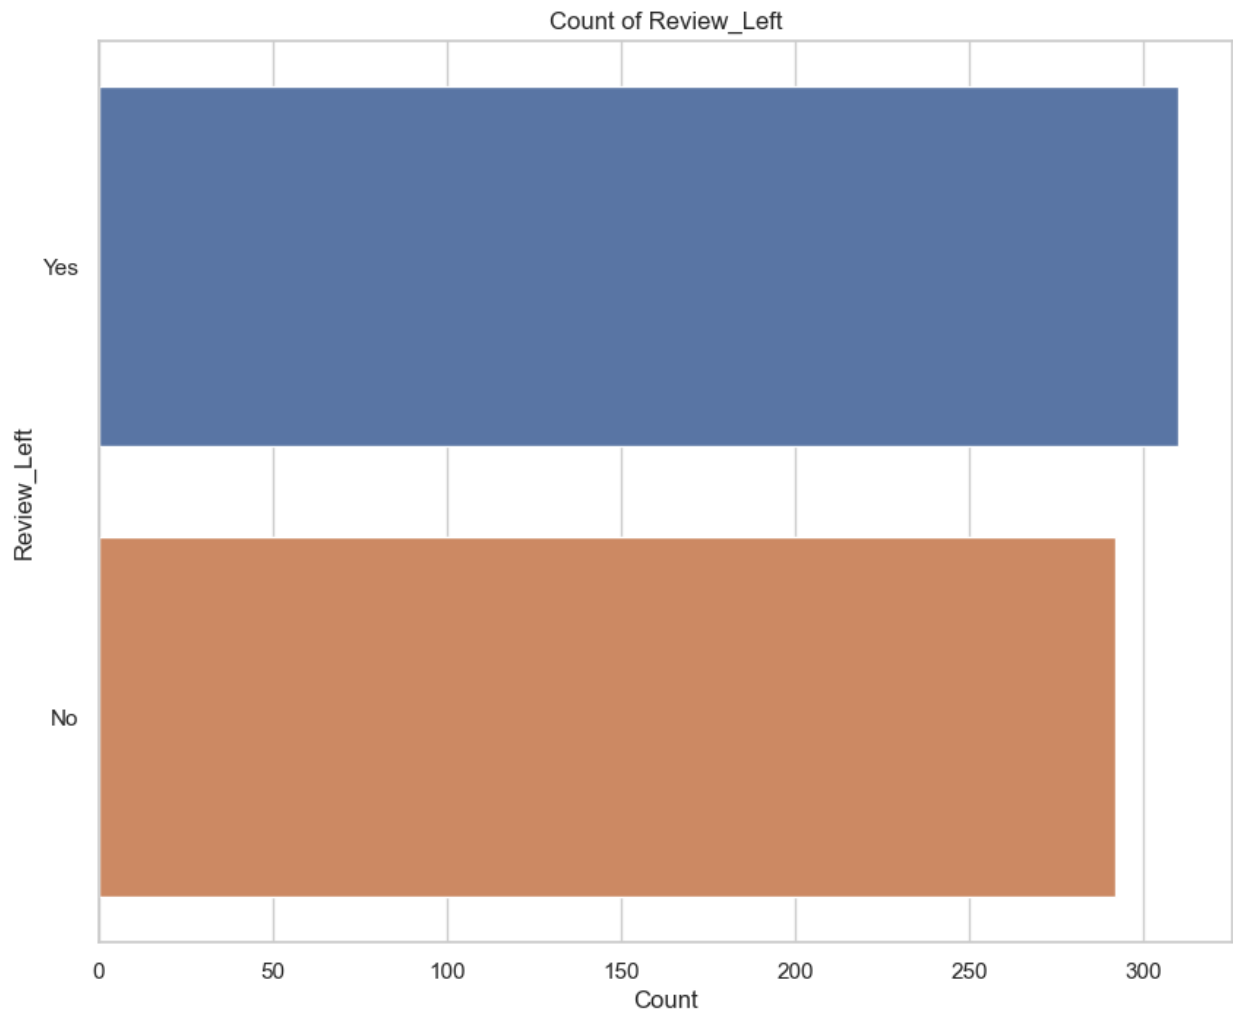
```
visualize_categorical(categorical['Review_Left'])
```

```
Review_Left
```

```
Yes    310
```

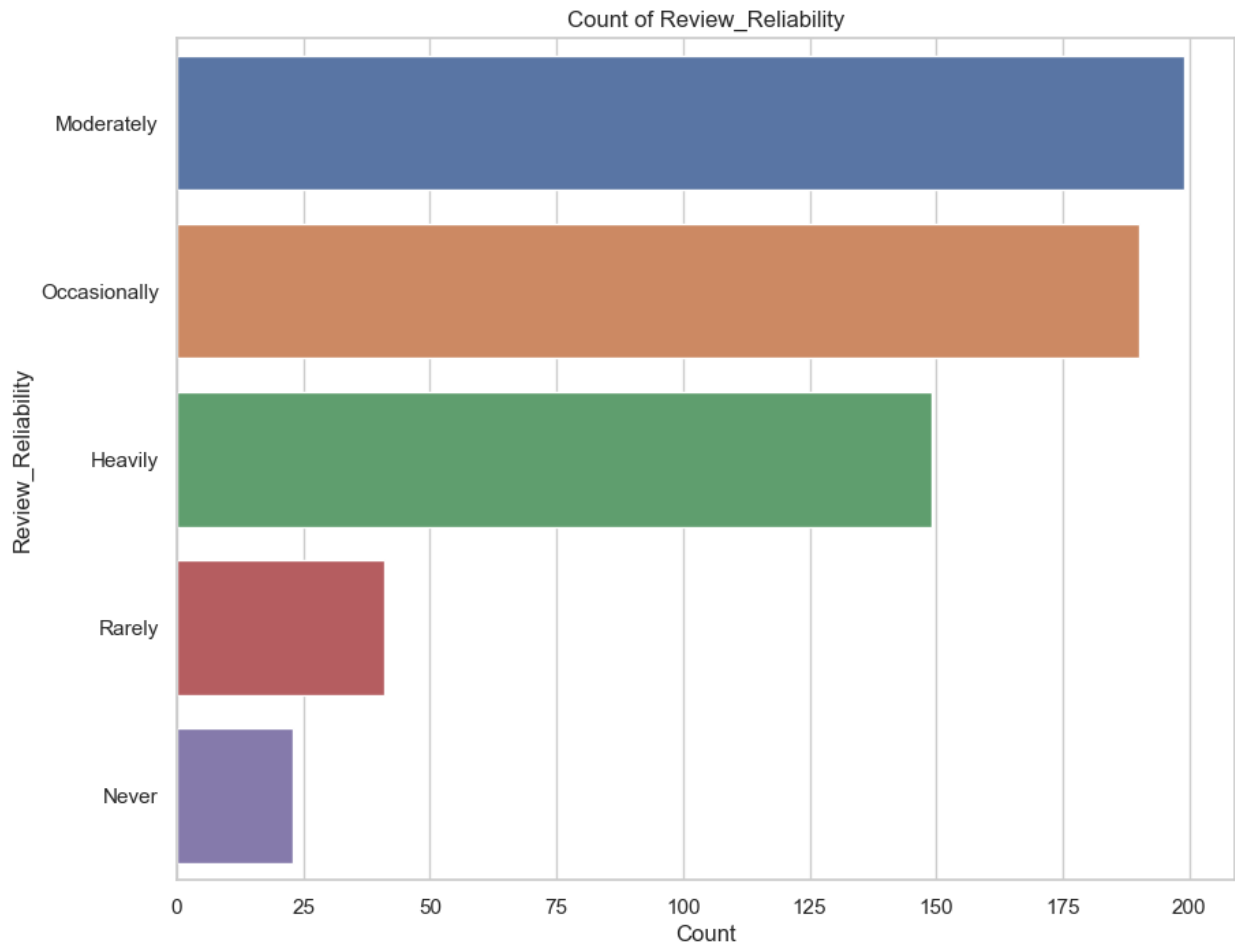
```
No     292
```

```
Name: count, dtype: int64
```



```
visualize_categorical(categorical['Review_Reliability'])
```

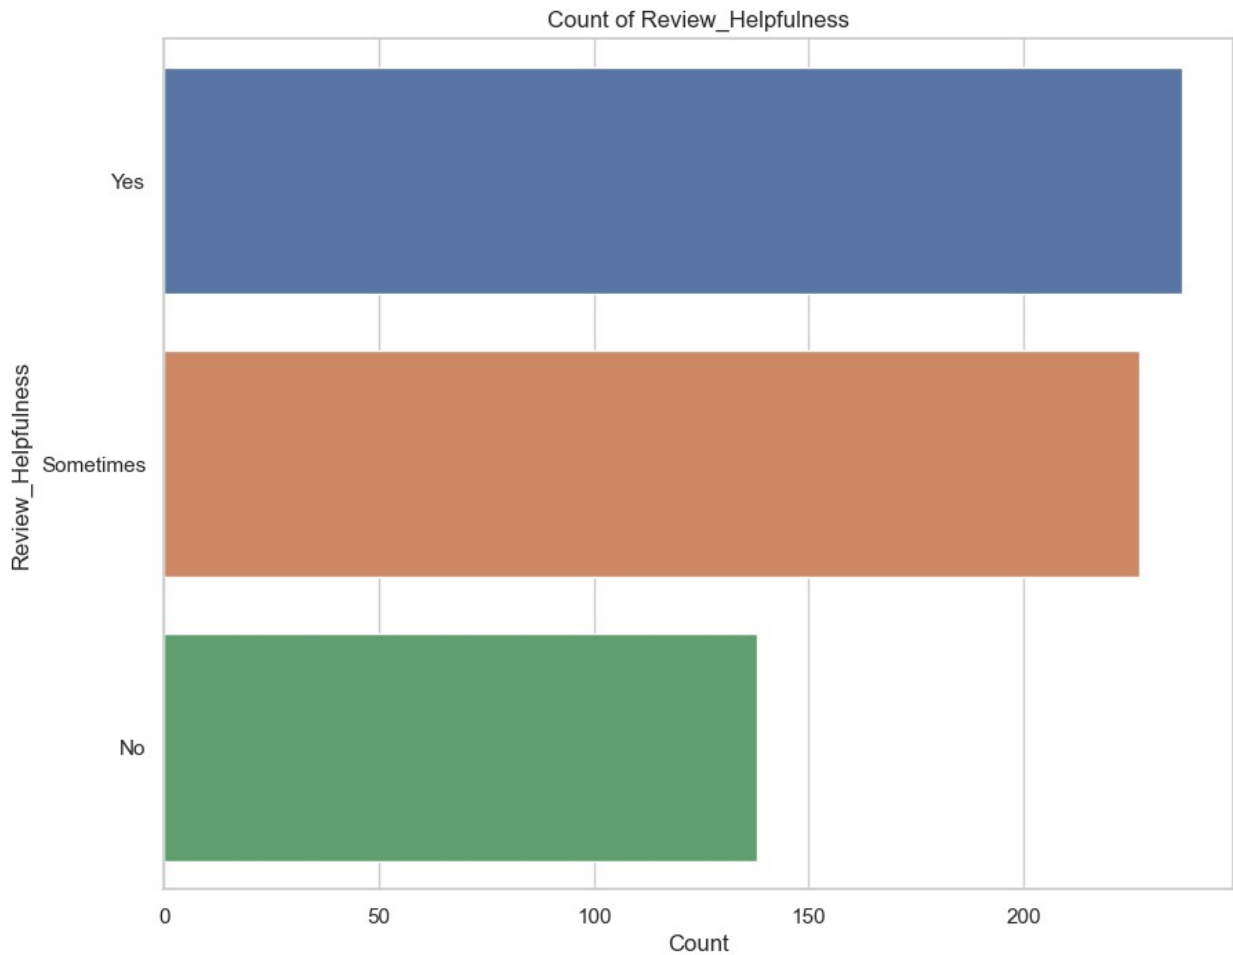
```
Review_Reliability
Moderately      199
Occasionally     190
Heavily          149
Rarely           41
Never            23
Name: count, dtype: int64
```



Most users find customer reviews helpful (Yes), with fewer users finding them only 'Sometimes' helpful and a small minority not finding them helpful ('No'). This emphasizes the importance of customer reviews in the shopping experience

```
visualize_categorical(categorical['Review_Helpfulness'])
```

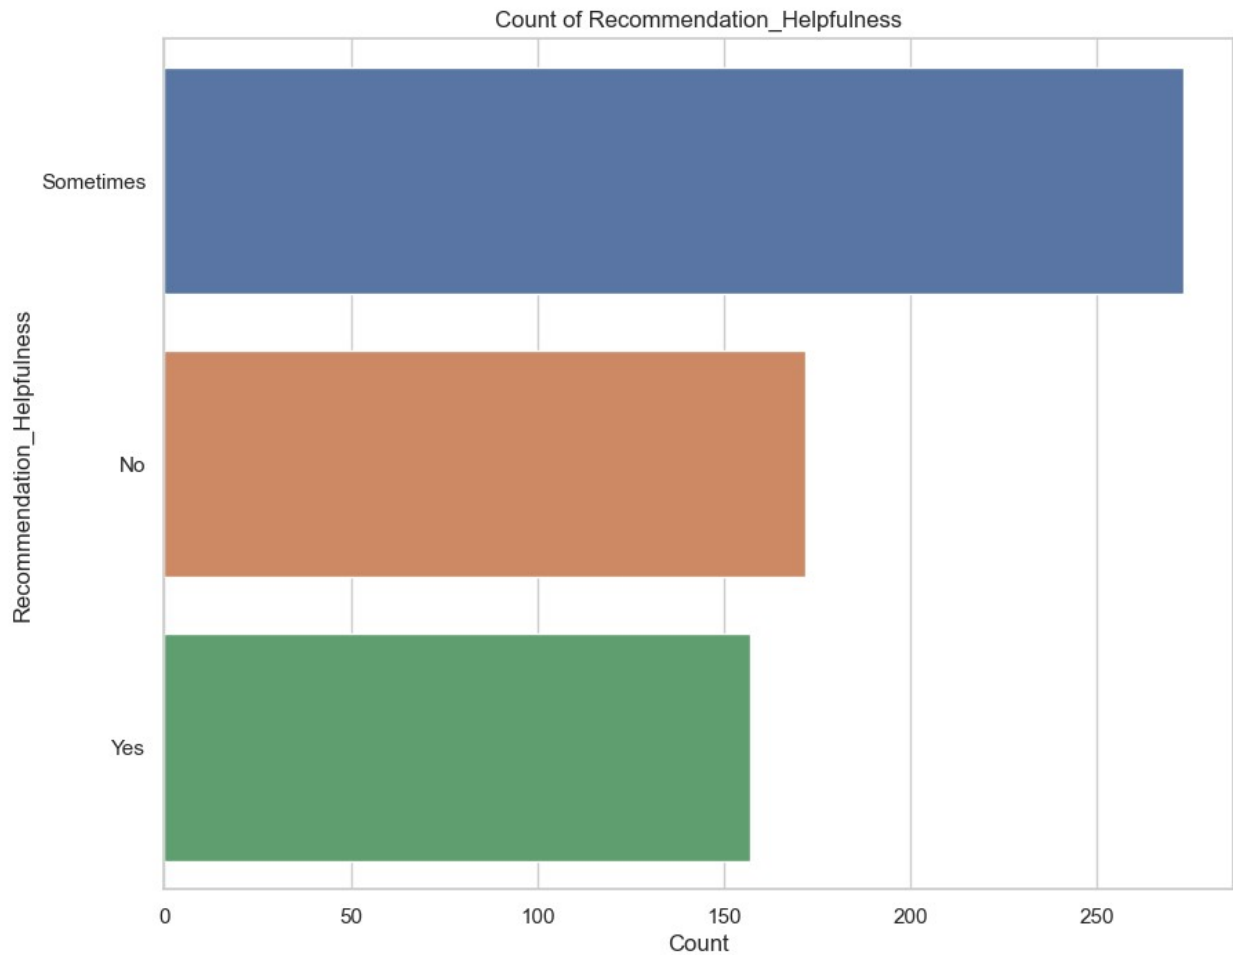
```
Review_Helpfulness
Yes                237
Sometimes          227
No                 138
Name: count, dtype: int64
```



Many customers find recommendations to be Sometimes helpful, indicating that while Amazon's recommendation system has an impact, it may not always be relevant or persuasive

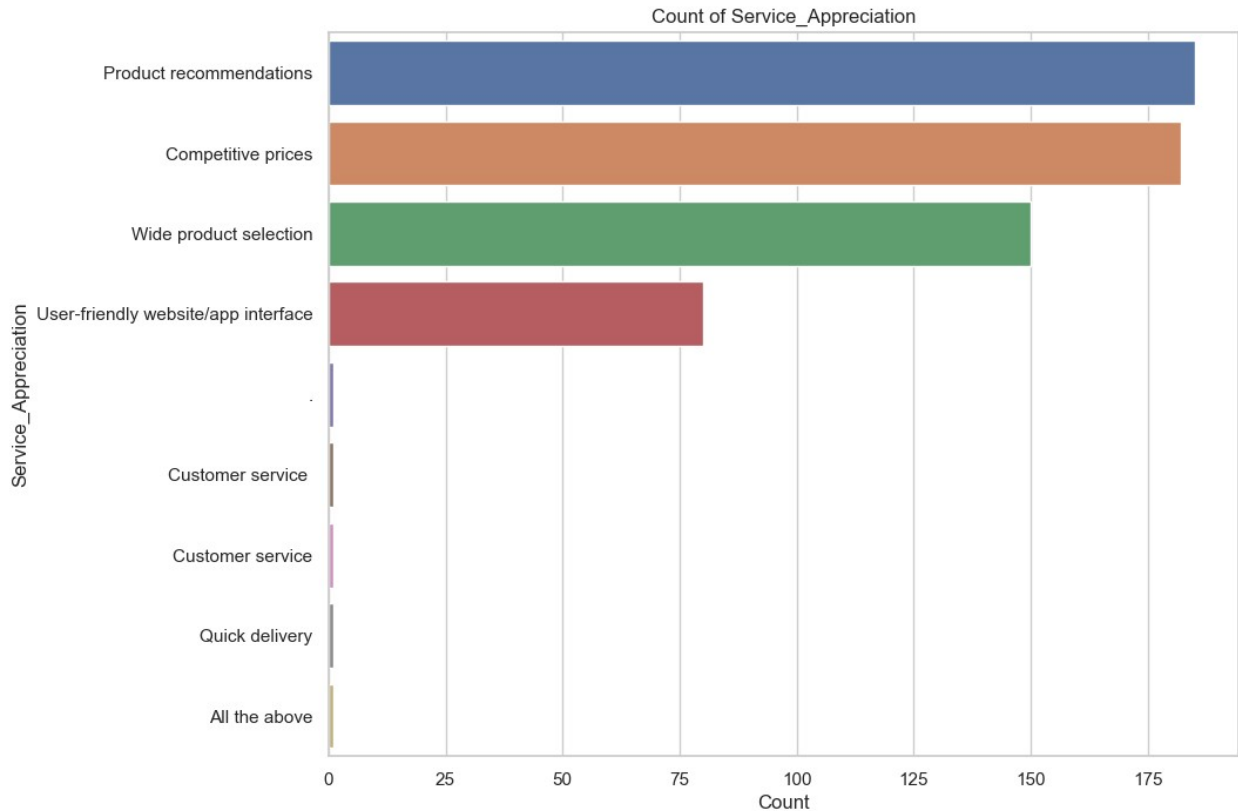
```
visualize_categorical(categorical['Recommendation_Helpfulness'])
```

```
Recommendation_Helpfulness
Sometimes    273
No           172
Yes          157
Name: count, dtype: int64
```



```
visualize_categorical(categorical['Service_Appreciation'])
```

```
Service_Appreciation
Product recommendations      185
Competitive prices          182
Wide product selection      150
User-friendly website/app interface  80
.                             1
Customer service            1
Customer service            1
Quick delivery              1
All the above               1
Name: count, dtype: int64
```



Product recommendations and Competitive prices are highly appreciated by users, followed closely by 'Wide product selection'. 'User-friendly website/app interface' and 'Quick delivery' are also valued but to a lesser extent. A small group appreciates 'All the above,' indicating overall satisfaction with multiple aspects of Amazon's services

```
visualize_categorical(categorical['Improvement_Areas'])
```

Improvement_Areas

Customer service responsiveness

217

Product quality and accuracy

159

Reducing packaging waste

133

Shipping speed and reliability

79

Quality of product is very poor according to the big offers

1

I don't have any problem with Amazon

1

User interface of app

1

Irrelevant product suggestions

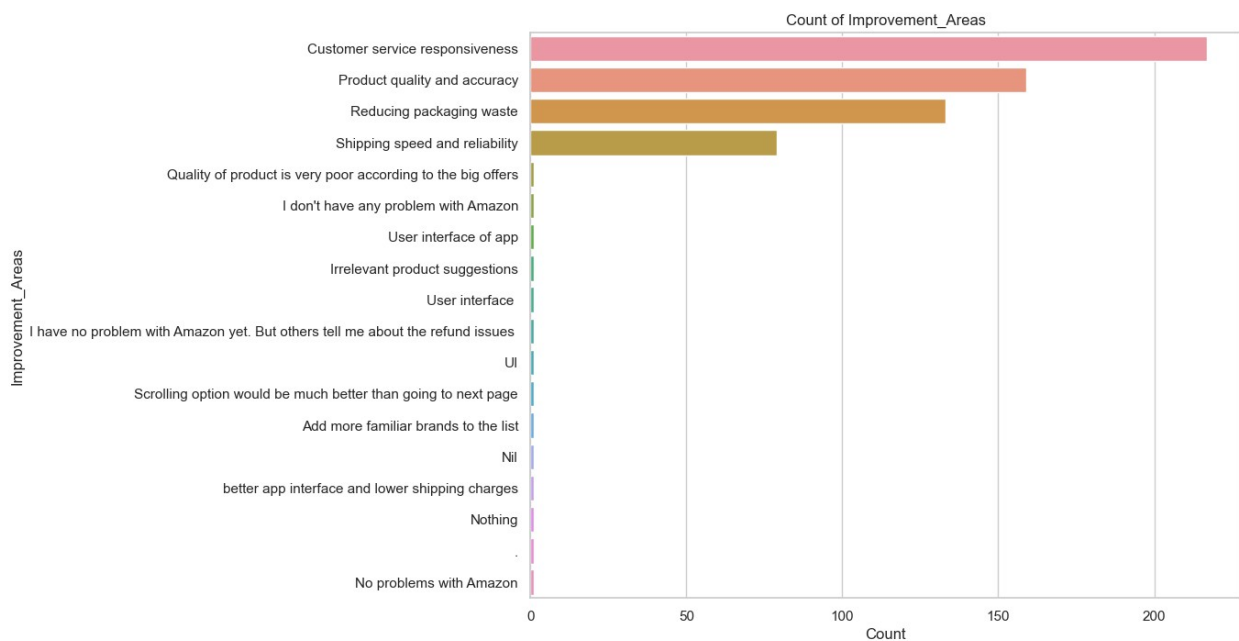
1

User interface

```

1
I have no problem with Amazon yet. But others tell me about the refund
issues      1
UI
1
Scrolling option would be much better than going to next page
1
Add more familiar brands to the list
1
Nil
1
better app interface and lower shipping charges
1
Nothing
1
.
1
No problems with Amazon
1
Name: count, dtype: int64

```



Customers appreciate Customer service responsiveness but also identify it as an area for improvement along with Product quality and accuracy and Shipping speed and reliability, suggesting these are important factors in customer satisfaction

Data Preprocessing

```

# Import necessary libraries
import pandas as pd
from sklearn.preprocessing import StandardScaler, LabelEncoder

```

```

from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

# Drop unnecessary columns
df_clean = df.drop(columns=['Timestamp'])

# Encode categorical variables using LabelEncoder
label_encoder = LabelEncoder()
for column in df_clean.select_dtypes(include=['object']).columns:
    df_clean[column] = label_encoder.fit_transform(df_clean[column])

```

Normalizing


```

# Normalize the data using StandardScaler
scaler = StandardScaler()
df_scaled = scaler.fit_transform(df_clean)

# Determine the optimal number of clusters using the elbow method
inertia = []
K = range(1, 11) # Test k from 1 to 10
for k in K:
    kmeans = KMeans(n_clusters=k, random_state=42)
    kmeans.fit(df_scaled)
    inertia.append(kmeans.inertia_)

# Plot the elbow graph to determine optimal k
plt.figure(figsize=(8, 6))
plt.plot(K, inertia, 'bo-', markersize=8)
plt.xlabel('Number of Clusters (k)')
plt.ylabel('Inertia')
plt.title('Elbow Method For Optimal k')
plt.show()

```

Model Evaluation: - Clustering

K-Means Clustering

K-Means is a widely used clustering technique that assigns customers into K clusters based on minimizing the distance between each customer and the centroid of the cluster. It's particularly useful for partitioning the dataset into a pre-specified number of clusters.

Steps:

- Normalize the dataset to ensure comparability across features.
- Use the Elbow Method to determine the optimal number of clusters.
- Apply the K-Means algorithm to segment the data.

```
# Choose k=3 (based on elbow method) and apply KMeans clustering
kmeans = KMeans(n_clusters=3, random_state=42)
df_clean['Cluster'] = kmeans.fit_predict(df_scaled)

# Analyze the characteristics of each cluster
cluster_analysis = df_clean.groupby('Cluster').mean()
print(cluster_analysis)
```

```

# Visualize the clusters using matplotlib
# Plotting based on 'age' and 'Purchase_Frequency' features
plt.figure(figsize=(10, 7))
colors = ['red', 'blue', 'green']
for cluster in range(3):
    plt.scatter(df_clean[df_clean['Cluster'] == cluster]['age'],
                df_clean[df_clean['Cluster'] == cluster]
                ['Purchase_Frequency'],
                c=colors[cluster], label=f'Cluster {cluster}')
plt.xlabel('Age')
plt.ylabel('Purchase Frequency')
plt.title('Customer Segments based on Age and Purchase Frequency')
plt.legend()
plt.show()

# Suggest marketing strategies based on cluster characteristics
for cluster in range(3):
    print(f"Cluster {cluster} Analysis:")

```

	age	Gender	Purchase_Frequency	Purchase_Categories
\ Cluster				
0	29.723810	0.847619	1.533333	13.352381
1	32.569444	0.881944	1.805556	6.524306
2	28.875598	0.497608	1.550239	12.315789

	Personalized_Recommendation_Frequency	Browsing_Frequency	\
Cluster			
0	0.571429	1.609524	
1	0.881944	0.770833	
2	0.760766	1.162679	

	Product_Search_Method	Search_Result_Exploration	\
Cluster			
0	1.800000	0.552381	
1	1.170139	0.711806	
2	1.196172	0.856459	

	Customer_Reviews_Importance	Add_to_Cart_Browsing	...	\
Cluster			...	
0	3.419048	1.076190	...	
1	2.809028	0.454861	...	
2	1.555024	1.559809	...	

Review_Helpfulness	Personalized_Recommendation_Frequency	\
--------------------	---------------------------------------	---

Cluster		
0	1.028571	3.657143
1	0.802083	2.701389
2	1.732057	2.215311

Recommendation_Helpfulness Rating_Accuracy
Shopping_Satisfaction \
Cluster

0	0.742857	3.619048
3.390476		
1	0.722222	2.673611
2.611111		
2	1.440191	2.196172
1.794258		

Service_Appreciation Improvement_Areas Hierarchical_Cluster
\
Cluster

0	5.666667	8.180952	2.361905
1	4.388889	6.170139	2.944444
2	5.760766	8.789474	1.473684

Cluster	DBSCAN_Cluster	Agglomerative_Cluster
0	-1.0	0.847619
1	-1.0	0.048611
2	-1.0	0.827751


[3 rows x 25 columns]


```

dendrogram(Z)
plt.title('Dendrogram')
plt.show()

# Extract clusters (decide number of clusters, e.g., 3)
clusters_hierarchical = fcluster(Z, 3, criterion='maxclust')
df_clean['Hierarchical_Cluster'] = clusters_hierarchical

```



DBSCAN (Density-Based Spatial Clustering of Applications with Noise)

DBSCAN identifies clusters based on density, which makes it effective for datasets with noise or uneven cluster shapes. It can also detect outliers (customers who don't belong to any cluster).

Steps:

- Select parameters based on the data distribution.
- Apply DBSCAN to cluster customers and identify noise (outliers).

```

# DBSCAN (Density-Based Spatial Clustering of Applications with Noise):

```

```


from sklearn.cluster import DBSCAN

# Apply DBSCAN
dbscan = DBSCAN(eps=0.5, min_samples=5)
clusters_dbscan = dbscan.fit_predict(df_scaled)

# Add cluster labels to the dataset
df_clean['DBSCAN_Cluster'] = clusters_dbscan

# Visualize DBSCAN clusters
plt.figure(figsize=(10, 7))
plt.scatter(df_clean['age'], df_clean['Purchase_Frequency'],
            c=clusters_dbscan, cmap='plasma')
plt.title('DBSCAN Clustering')
plt.xlabel('Age')
plt.ylabel('Purchase Frequency')
plt.show()

```



Analysis:

- Cluster 0: Consistent spenders with high purchase frequency.

- Cluster 1: Irregular, low-frequency buyers.
- Outliers: Customers with unusual purchasing patterns, possibly anomalous behavior or new custom.

PCA (Principal Component Analysis)

PCA is not a clustering algorithm but a dimensionality reduction technique that helps visualize high-dimensional data by projecting it onto two principal components. This helps in visualizing the clusters formed by K-Means or other clustering algorithms.

Steps:

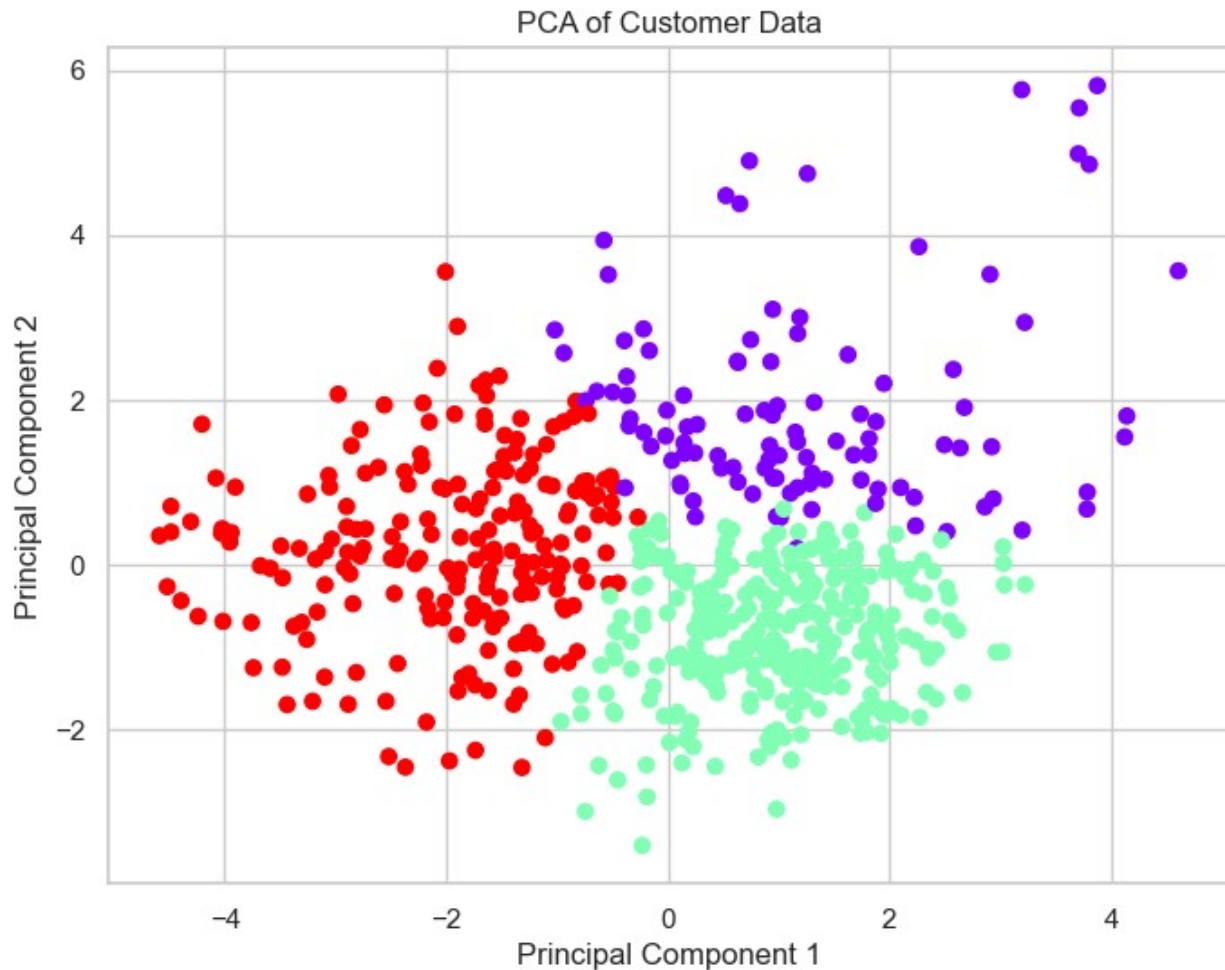
- Reduce the dimensionality of the dataset to two components.
- Visualize clusters in a 2D space.

```
# PCA (Principal Component Analysis):

from sklearn.decomposition import PCA

# Apply PCA to reduce dimensions to 2
pca = PCA(n_components=2)
df_pca = pca.fit_transform(df_scaled)

# Visualize the data in 2D after PCA
plt.figure(figsize=(8, 6))
plt.scatter(df_pca[:, 0], df_pca[:, 1], c=kmeans.labels_,
            cmap='rainbow')
plt.title('PCA of Customer Data')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.show()
```



Analysis:

- The PCA plot shows how well-separated the clusters are in a two-dimensional space.
- This is a useful tool for validating clustering results and identifying overlaps between clusters.

Agglomerative Clustering:

It is a bottom-up approach to hierarchical clustering, where each customer starts as its own cluster, and clusters are merged based on similarity until one single cluster remains.

Steps:

- Use Ward's method to minimize within-cluster variance.
- Plot a dendrogram to visualize the hierarchical structure of clusters.
- Cut the dendrogram to assign customers to a specific number of clusters.

```
# Agglomerative Clustering:
```

```
from sklearn.cluster import AgglomerativeClustering
```




```

# Apply Agglomerative Clustering
agglo = AgglomerativeClustering(n_clusters=3)
clusters_agglo = agglo.fit_predict(df_scaled)

# Add cluster labels to the dataset
df_clean['Agglomerative_Cluster'] = clusters_agglo

# Visualize Agglomerative clusters
plt.figure(figsize=(10, 7))
plt.scatter(df_clean['age'], df_clean['Purchase_Frequency'],
            c=clusters_agglo, cmap='cool')
plt.title('Agglomerative Clustering')
plt.xlabel('Age')
plt.ylabel('Purchase Frequency')
plt.show()

```



Analysis:

- Cluster 1: Customers who frequently purchase low-cost products.
- Cluster 2: High-value but infrequent buyers.
- Cluster 3: Moderate frequency and medium-value buyers.

Model Comparisons

- **K-Means Clustering:** Offers efficient segmentation based on spending and purchase frequency but requires a predefined number of clusters.
- **Agglomerative Clustering:** Provides a hierarchical view of customer relationships, allowing flexibility in the number of clusters. The dendrogram helps visualize the clustering structure.
- **DBSCAN:** Suitable for detecting noise (outliers) and identifying clusters of irregular shapes without predefining the number of clusters. Ideal for datasets with non-uniform cluster shapes.
- **PCA:** Helps in visualizing and validating clusters in reduced dimensions, offering insight into the relationships between customers in two principal components.

Each clustering technique has its strengths and weaknesses. By combining them, businesses can better understand their customer base and develop effective marketing strategies tailored to distinct customer segments. For example, high-frequency buyers can receive loyalty programs, while outliers may need special offers to encourage engagement.

Other Models Evaluation

```
# import libraries
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from pandas.plotting import scatter_matrix
from sklearn.neighbors import LocalOutlierFactor
from sklearn.preprocessing import OrdinalEncoder, LabelEncoder,
StandardScaler
from sklearn.feature_selection import SelectKBest, chi2
from sklearn.model_selection import train_test_split, KFold,
cross_val_score, GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.pipeline import Pipeline
from sklearn.ensemble import AdaBoostClassifier,
GradientBoostingClassifier, RandomForestClassifier,
ExtraTreesClassifier
from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score

# split input and output variable
x = df.drop(['Gender'], axis=1)
y = df['Gender']

# convert to string
x = x.astype(str)
```

Catagorical Encoding

```
# prepare input data
def prepare_inputs(x_train, x_test):
    oe = OrdinalEncoder(handle_unknown='use_encoded_value',
unknown_value=-1)
    oe.fit(x_train)
    x_train_enc = oe.transform(x_train)
    x_test_enc = oe.transform(x_test)
    return x_train_enc, x_test_enc
```

Target Encoding

```
# prepare target
def prepare_targets(y_train, y_test):
    le = LabelEncoder()
    le.fit(y_train)
    y_train_enc = le.transform(y_train)
    y_test_enc = le.transform(y_test)
    return y_train_enc, y_test_enc
```

Feature Scaling

```
# feature selection
def select_features(x_train, y_train, x_test):
    fs = SelectKBest(score_func=chi2, k='all')
    fs.fit(x_train, y_train)
    x_train_fs = fs.transform(x_train)
    x_test_fs = fs.transform(x_test)
    return x_train_fs, x_test_fs, fs

# split the dataset
x_train, x_test, y_train, y_test = train_test_split(x, y,
test_size=0.2, random_state=23)
x_train.shape, x_test.shape, y_train.shape, y_test.shape

((481, 22), (121, 22), (481,), (121,))

# prepare input data
x_train_enc, x_test_enc = prepare_inputs(x_train, x_test)


# prepare output data
y_train_enc, y_test_enc = prepare_targets(y_train, y_test)

# feature selection
x_train_fs, x_test_fs, fs = select_features(x_train_enc, y_train_enc,
x_test_enc)
```

```
# what are scores for the features
for i in range(len(fs.scores_)):
    print('Feature %d: %f' % (i, fs.scores_[i]))
```

```
Feature 0: 673.072125
Feature 1: 30.231644
Feature 2: 0.982145
Feature 3: 196.669813
Feature 4: 2.882849
Feature 5: 1.639115
Feature 6: 5.355460
Feature 7: 1.663367
Feature 8: 13.991366
Feature 9: 9.591061
Feature 10: 4.185675
Feature 11: 1.355832
Feature 12: 2.441824
Feature 13: 0.706066
Feature 14: 7.305193
Feature 15: 10.537447
Feature 16: 7.709323
Feature 17: 4.472732
Feature 18: 3.420467
Feature 19: 10.136995
Feature 20: 7.125628
Feature 21: 6.940209
```

```
# plot the scores
plt.bar([i for i in range(len(fs.scores_))], fs.scores_)
plt.show()
```



Algorithms


```
# spot check algorithms
models = []
models.append(("LR", LogisticRegression(solver="lbfgs",
max_iter=1000)))
models.append(("LDA", LinearDiscriminantAnalysis()))
models.append(("DT", DecisionTreeClassifier()))
models.append(("KNN", KNeighborsClassifier()))
models.append(("NB", GaussianNB()))
models.append(("SVM", SVC()))

# evaluate each model in turn
results = []
names = []
for name, model in models:
    cv = KFold(n_splits=10, random_state=None)
    scores = cross_val_score(model, x_train_enc, y_train_enc,
scoring="accuracy", cv=cv)
    names.append(name)
    results.append(scores)
    print("%s %.2f (%.2f)" % (name, scores.mean(), scores.std()))

LR 0.59 (0.05)
LDA 0.60 (0.04)
DT 0.51 (0.05)
```

```
KNN 0.58 (0.05)
NB 0.50 (0.06)
SVM 0.61 (0.04)
```

```
# compare algorithm
fig = plt.figure()
fig.suptitle("Algorithm Comparison")
ax = fig.add_subplot(111)
plt.boxplot(results)
ax.set_xticklabels(names)
plt.show()
```



Hyperparameter Tuning with grid search

```
# tuned with svm
c_values = [0.1, 0.3, 0.5, 0.7, 0.9, 1.0, 1.3, 1.5, 1.7, 2.0]
kernel_values = ['linear', 'poly', 'rbf', 'sigmoid']
param_grid = dict(C=c_values, kernel=kernel_values)
model = SVC()
cv = KFold(n_splits=10, random_state=None)
grid = GridSearchCV(estimator=model, param_grid=param_grid,
                    scoring="accuracy", cv=cv)
```

```

grid_result = grid.fit(x_train_enc, y_train_enc)
print("Best: %.3f using %r" % (grid_result.best_score_,
grid_result.best_params_))
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, std, param in zip(means, stds, params):
    print("%.3f (%.3f) with %r" % (mean, std, param))

```

```

Best: 0.607 using {'C': 0.1, 'kernel': 'poly'}
0.605 (0.043) with {'C': 0.1, 'kernel': 'linear'}
0.607 (0.044) with {'C': 0.1, 'kernel': 'poly'}
0.607 (0.044) with {'C': 0.1, 'kernel': 'rbf'}
0.607 (0.044) with {'C': 0.1, 'kernel': 'sigmoid'}
0.605 (0.043) with {'C': 0.3, 'kernel': 'linear'}
0.607 (0.044) with {'C': 0.3, 'kernel': 'poly'}
0.607 (0.044) with {'C': 0.3, 'kernel': 'rbf'}
0.551 (0.068) with {'C': 0.3, 'kernel': 'sigmoid'}
0.605 (0.043) with {'C': 0.5, 'kernel': 'linear'}
0.607 (0.044) with {'C': 0.5, 'kernel': 'poly'}
0.607 (0.044) with {'C': 0.5, 'kernel': 'rbf'}
0.528 (0.082) with {'C': 0.5, 'kernel': 'sigmoid'}
0.605 (0.043) with {'C': 0.7, 'kernel': 'linear'}
0.607 (0.044) with {'C': 0.7, 'kernel': 'poly'}
0.607 (0.044) with {'C': 0.7, 'kernel': 'rbf'}
0.526 (0.089) with {'C': 0.7, 'kernel': 'sigmoid'}
0.605 (0.043) with {'C': 0.9, 'kernel': 'linear'}
0.607 (0.044) with {'C': 0.9, 'kernel': 'poly'}
0.607 (0.044) with {'C': 0.9, 'kernel': 'rbf'}
0.519 (0.090) with {'C': 0.9, 'kernel': 'sigmoid'}
0.605 (0.043) with {'C': 1.0, 'kernel': 'linear'}
0.607 (0.044) with {'C': 1.0, 'kernel': 'poly'}
0.607 (0.044) with {'C': 1.0, 'kernel': 'rbf'}
0.519 (0.090) with {'C': 1.0, 'kernel': 'sigmoid'}
0.605 (0.043) with {'C': 1.3, 'kernel': 'linear'}
0.607 (0.044) with {'C': 1.3, 'kernel': 'poly'}
0.607 (0.044) with {'C': 1.3, 'kernel': 'rbf'}
0.513 (0.086) with {'C': 1.3, 'kernel': 'sigmoid'}
0.605 (0.043) with {'C': 1.5, 'kernel': 'linear'}
0.607 (0.044) with {'C': 1.5, 'kernel': 'poly'}
0.607 (0.044) with {'C': 1.5, 'kernel': 'rbf'}
0.513 (0.086) with {'C': 1.5, 'kernel': 'sigmoid'}
0.605 (0.043) with {'C': 1.7, 'kernel': 'linear'}
0.607 (0.044) with {'C': 1.7, 'kernel': 'poly'}
0.607 (0.044) with {'C': 1.7, 'kernel': 'rbf'}
0.507 (0.088) with {'C': 1.7, 'kernel': 'sigmoid'}
0.605 (0.043) with {'C': 2.0, 'kernel': 'linear'}
0.607 (0.044) with {'C': 2.0, 'kernel': 'poly'}
0.607 (0.044) with {'C': 2.0, 'kernel': 'rbf'}
0.511 (0.087) with {'C': 2.0, 'kernel': 'sigmoid'}

```

```

# standardize with basic algorithms
pipelines = []
pipelines.append(("scalerLR", Pipeline([("scaler", StandardScaler()),
("LR", LogisticRegression(solver="lbfgs", max_iter=1000))])))
pipelines.append(("scalerLDA", Pipeline([("scaler", StandardScaler()),
("LDA", LinearDiscriminantAnalysis())])))
pipelines.append(("scalerKNN", Pipeline([("scaler", StandardScaler()),
("KNN", KNeighborsClassifier())])))
pipelines.append(("scalerNB", Pipeline([("scaler", StandardScaler()),
("NB", GaussianNB())])))
pipelines.append(("scalerCART", Pipeline([("scaler",
StandardScaler()), ("CART", DecisionTreeClassifier())])))
pipelines.append(("scalerSVM", Pipeline([("scaler", StandardScaler()),
("SVM", SVC())])))


# evaluate algorithm with standardize
names = []
results = []
for name, model in pipelines:
    cv = KFold(n_splits=10, random_state=None)
    scores = cross_val_score(model, x_train_enc, y_train_enc,
scoring="accuracy", cv=cv)
    names.append(name)
    results.append(scores)
    print("%s %.3f (%.3f)" % (name, scores.mean(), scores.std()))

scalerLR 0.601 (0.054)
scalerLDA 0.599 (0.045)
scalerKNN 0.576 (0.051)
scalerNB 0.501 (0.061)
scalerCART 0.485 (0.069)
scalerSVM 0.603 (0.045)

# compare algorithm
fig = plt.figure()
fig.suptitle("Scaler Algorithm Comparison")
ax = fig.add_subplot(111)
plt.boxplot(results)
ax.set_xticklabels(names)
plt.show()

```


Scaler Algorithm Comparison



Ensamble Method

```
# tuned with svm standardize
scaler = StandardScaler()
rescaled_x = scaler.fit_transform(x_train_enc)
c_values = [0.1, 0.3, 0.5, 0.7, 0.9, 1.0, 1.3, 1.5, 1.7, 2.0]
kernel_values = ['linear', 'poly', 'rbf', 'sigmoid']
param_grid = dict(C=c_values, kernel=kernel_values)
model = SVC()
cv = KFold(n_splits=10, random_state=None)
grid = GridSearchCV(estimator=model, param_grid=param_grid,
                    scoring="accuracy", cv=cv)
grid_result = grid.fit(rescaled_x, y_train_enc)
print("Best: %.3f using %r" % (grid_result.best_score_,
                              grid_result.best_params_))
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, std, param in zip(means, stds, params):
    print("%.3f (%.3f) with %r" % (mean, std, param))

Best: 0.609 using {'C': 1.7, 'kernel': 'sigmoid'}
0.605 (0.043) with {'C': 0.1, 'kernel': 'linear'}
```

```
0.607 (0.044) with {'C': 0.1, 'kernel': 'poly'}
0.607 (0.044) with {'C': 0.1, 'kernel': 'rbf'}
0.607 (0.044) with {'C': 0.1, 'kernel': 'sigmoid'}
0.605 (0.043) with {'C': 0.3, 'kernel': 'linear'}
0.605 (0.043) with {'C': 0.3, 'kernel': 'poly'}
0.607 (0.044) with {'C': 0.3, 'kernel': 'rbf'}
0.607 (0.044) with {'C': 0.3, 'kernel': 'sigmoid'}
0.605 (0.043) with {'C': 0.5, 'kernel': 'linear'}
0.607 (0.040) with {'C': 0.5, 'kernel': 'poly'}
0.607 (0.044) with {'C': 0.5, 'kernel': 'rbf'}
0.607 (0.044) with {'C': 0.5, 'kernel': 'sigmoid'}
0.605 (0.043) with {'C': 0.7, 'kernel': 'linear'}
0.603 (0.045) with {'C': 0.7, 'kernel': 'poly'}
0.607 (0.044) with {'C': 0.7, 'kernel': 'rbf'}
0.607 (0.044) with {'C': 0.7, 'kernel': 'sigmoid'}
0.605 (0.043) with {'C': 0.9, 'kernel': 'linear'}
0.607 (0.044) with {'C': 0.9, 'kernel': 'poly'}
0.607 (0.044) with {'C': 0.9, 'kernel': 'rbf'}
0.607 (0.044) with {'C': 0.9, 'kernel': 'sigmoid'}
0.605 (0.043) with {'C': 1.0, 'kernel': 'linear'}
0.605 (0.047) with {'C': 1.0, 'kernel': 'poly'}
0.603 (0.045) with {'C': 1.0, 'kernel': 'rbf'}
0.607 (0.044) with {'C': 1.0, 'kernel': 'sigmoid'}
0.605 (0.043) with {'C': 1.3, 'kernel': 'linear'}
0.601 (0.048) with {'C': 1.3, 'kernel': 'poly'}
0.592 (0.047) with {'C': 1.3, 'kernel': 'rbf'}
0.607 (0.044) with {'C': 1.3, 'kernel': 'sigmoid'}
0.605 (0.043) with {'C': 1.5, 'kernel': 'linear'}
0.603 (0.048) with {'C': 1.5, 'kernel': 'poly'}
0.597 (0.046) with {'C': 1.5, 'kernel': 'rbf'}
0.607 (0.042) with {'C': 1.5, 'kernel': 'sigmoid'}
0.605 (0.043) with {'C': 1.7, 'kernel': 'linear'}
0.601 (0.048) with {'C': 1.7, 'kernel': 'poly'}
0.586 (0.052) with {'C': 1.7, 'kernel': 'rbf'}
0.609 (0.045) with {'C': 1.7, 'kernel': 'sigmoid'}
0.605 (0.043) with {'C': 2.0, 'kernel': 'linear'}
0.599 (0.049) with {'C': 2.0, 'kernel': 'poly'}
0.590 (0.053) with {'C': 2.0, 'kernel': 'rbf'}
0.607 (0.042) with {'C': 2.0, 'kernel': 'sigmoid'}
```

```
# ensemble methods
```

```
ensembles = []
ensembles.append(('AB', AdaBoostClassifier()))
ensembles.append(('GBM', GradientBoostingClassifier()))
ensembles.append(('RF', RandomForestClassifier()))
ensembles.append(('ET', ExtraTreesClassifier()))
```

```
# evaluate each model with ensemble
```

```
results = []
names = []
```

```

for name, model in ensembles:
    kfold = KFold(n_splits=10, random_state=None)
    cv_results = cross_val_score(model, x_train_enc, y_train_enc,
cv=kfold, scoring="accuracy")
    results.append(cv_results)
    names.append(name)
    print("%s: %.3f (%.3f)" % (name, cv_results.mean(),
cv_results.std()))

```

AB: 0.497 (0.088)

GBM: 0.557 (0.060)

RF: 0.599 (0.051)

ET: 0.615 (0.050)

Compare Algorithms

```
fig = plt.figure()
```

```
fig.suptitle('Ensemble Algorithm Comparison')
```


```
ax = fig.add_subplot(111)
```

```
plt.boxplot(results)
```

```
ax.set_xticklabels(names)
```

```
plt.show()
```

Ensemble Algorithm Comparison



```

# tuned with ETC
estimator = [10, 20, 50, 100, 200, 500, 1000]
criterion= ['gini', 'entropy']
param_grid = dict(n_estimators=estimator, criterion=criterion)
model = ExtraTreesClassifier()
cv = KFold(n_splits=10, random_state=None)
grid = GridSearchCV(estimator=model, param_grid=param_grid,
scoring="accuracy", cv=cv)
grid_result = grid.fit(x_train_enc, y_train_enc)
print("Best: %.3f using %r" % (grid_result.best_score_,
grid_result.best_params_))
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, std, param in zip(means, stds, params):
    print("%.3f (%.3f) with %r" % (mean, std, param))

Best: 0.605 using {'criterion': 'gini', 'n_estimators': 1000}
0.557 (0.058) with {'criterion': 'gini', 'n_estimators': 10}
0.586 (0.053) with {'criterion': 'gini', 'n_estimators': 20}
0.574 (0.049) with {'criterion': 'gini', 'n_estimators': 50}
0.599 (0.059) with {'criterion': 'gini', 'n_estimators': 100}
0.603 (0.047) with {'criterion': 'gini', 'n_estimators': 200}
0.603 (0.050) with {'criterion': 'gini', 'n_estimators': 500}
0.605 (0.051) with {'criterion': 'gini', 'n_estimators': 1000}
0.561 (0.063) with {'criterion': 'entropy', 'n_estimators': 10}
0.578 (0.057) with {'criterion': 'entropy', 'n_estimators': 20}
0.588 (0.053) with {'criterion': 'entropy', 'n_estimators': 50}
0.605 (0.054) with {'criterion': 'entropy', 'n_estimators': 100}
0.605 (0.053) with {'criterion': 'entropy', 'n_estimators': 200}
0.601 (0.055) with {'criterion': 'entropy', 'n_estimators': 500}
0.603 (0.055) with {'criterion': 'entropy', 'n_estimators': 1000}

```

```

# ensembles with standardize
ensembles = []
ensembles.append(("scalerRF", Pipeline([("scaler", StandardScaler()),
("RF", RandomForestClassifier())])))
ensembles.append(("scalerGBM", Pipeline([("scaler", StandardScaler()),
("GBM", GradientBoostingClassifier())])))
ensembles.append(("scalerET", Pipeline([("scaler", StandardScaler()),
("ET", ExtraTreesClassifier())])))
ensembles.append(("scalerAB", Pipeline([("scaler", StandardScaler()),
("AB", AdaBoostClassifier())])))

```

```

# evaluate each model with ensemble
results = []
names = []
for name, model in ensembles:
    kfold = KFold(n_splits=10, random_state=None)
    cv_results = cross_val_score(model, x_train_enc, y_train_enc,

```

```


cv=kfold, scoring="accuracy")
    results.append(cv_results)
    names.append(name)
    print("%s: %.3f (%.3f)" % (name, cv_results.mean(),
cv_results.std()))

scalerRF: 0.588 (0.053)
scalerGBM: 0.555 (0.064)
scalerET: 0.590 (0.054)
scalerAB: 0.497 (0.088)

# Compare Algorithms
fig = plt.figure()
fig.suptitle('Ensemble Algorithm Comparison with Standardization')
ax = fig.add_subplot(111)
plt.boxplot(results)
ax.set_xticklabels(names)
plt.show()

```

Ensemble Algorithm Comparison with Standardization



```

# tuned with RF
scaler = StandardScaler()
rescaled_x = scaler.fit_transform(x_train_enc)

```

```

estimator = [10, 20, 50, 100, 200, 500, 1000]
criterion= ['gini', 'entropy']
param_grid = dict(n_estimators=estimator, criterion=criterion)#
prepare the model
model = ExtraTreesClassifier(n_estimators=50, criterion="entropy")
model.fit(x_train_enc, y_train_enc)

model = ExtraTreesClassifier(random_state=42)
cv = KFold(n_splits=10, random_state=None)
grid = GridSearchCV(estimator=model, param_grid=param_grid,
                    scoring="accuracy", cv=cv)
grid_result = grid.fit(rescaled_x, y_train_enc)
print("Best: %.3f using %r" % (grid_result.best_score_,
                               grid_result.best_params_))
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, std, param in zip(means, stds, params):
    print("%.3f (%.3f) with %r" % (mean, std, param))

Best: 0.607 using {'criterion': 'entropy', 'n_estimators': 100}
0.549 (0.061) with {'criterion': 'gini', 'n_estimators': 10}
0.576 (0.052) with {'criterion': 'gini', 'n_estimators': 20}
0.588 (0.051) with {'criterion': 'gini', 'n_estimators': 50}
0.597 (0.064) with {'criterion': 'gini', 'n_estimators': 100}
0.592 (0.062) with {'criterion': 'gini', 'n_estimators': 200}
0.599 (0.056) with {'criterion': 'gini', 'n_estimators': 500}
0.599 (0.053) with {'criterion': 'gini', 'n_estimators': 1000}
0.580 (0.049) with {'criterion': 'entropy', 'n_estimators': 10}
0.594 (0.059) with {'criterion': 'entropy', 'n_estimators': 20}
0.599 (0.056) with {'criterion': 'entropy', 'n_estimators': 50}
0.607 (0.049) with {'criterion': 'entropy', 'n_estimators': 100}
0.603 (0.051) with {'criterion': 'entropy', 'n_estimators': 200}
0.603 (0.057) with {'criterion': 'entropy', 'n_estimators': 500}
0.607 (0.055) with {'criterion': 'entropy', 'n_estimators': 1000}

```

Finalize the model

```

# prepare the model
model = ExtraTreesClassifier(n_estimators=50, criterion="entropy")
model.fit(x_train_enc, y_train_enc)

ExtraTreesClassifier(criterion='entropy', n_estimators=50)

# estimate accuracy on validation dataset
predictions = model.predict(x_test_enc)
print(accuracy_score(y_test_enc, predictions))
print(confusion_matrix(y_test_enc, predictions))
print(classification_report(y_test_enc, predictions))

```

0.4793388429752066

```
[[57  2  0  1]
 [33  1  0  1]
 [ 4  0  0  0]
 [21  1  0  0]]
```

	precision	recall	f1-score	support
0	0.50	0.95	0.65	60
1	0.25	0.03	0.05	35
2	0.00	0.00	0.00	4
3	0.00	0.00	0.00	22
accuracy			0.48	121
macro avg	0.19	0.24	0.18	121
weighted avg	0.32	0.48	0.34	121

Recommendations:

Targeted Marketing Strategies Based on Cluster Analysis:

Frequent Buyers, Moderate Spending:

- Loyalty Programs: Offer points or discounts for frequent purchases.
- Cross-Selling: Recommend complementary products to increase average order value.

Low-Frequency, High-Spending Buyers:

- Exclusive Offers: Provide limited-time deals on premium products to encourage more frequent purchases.
- VIP Treatment: Send personalized offers or early access to sales.

Moderate Buyers, Balanced Spending:

- Personalized Discounts: Offer tailored promotions based on past purchases.
- Product Bundling: Encourage bundled purchases to maximize value.

These strategies aim to drive engagement, retention, and increased spending across each segment.

References

- Mighty Itama Itama, PhD. (2024). Machine Learning using Python.
<https://amightyo.quarto.pub/machine-learning-using-python/>